

Chapter 2

Operating System Overview

EIGHTH EDITION

BY WILLIAM STALLINGS

Operating System

A program that controls the execution of application programs

An interface between applications and hardware

Main objectives of an OS:

- convenience
- efficiency
- ability to evolve

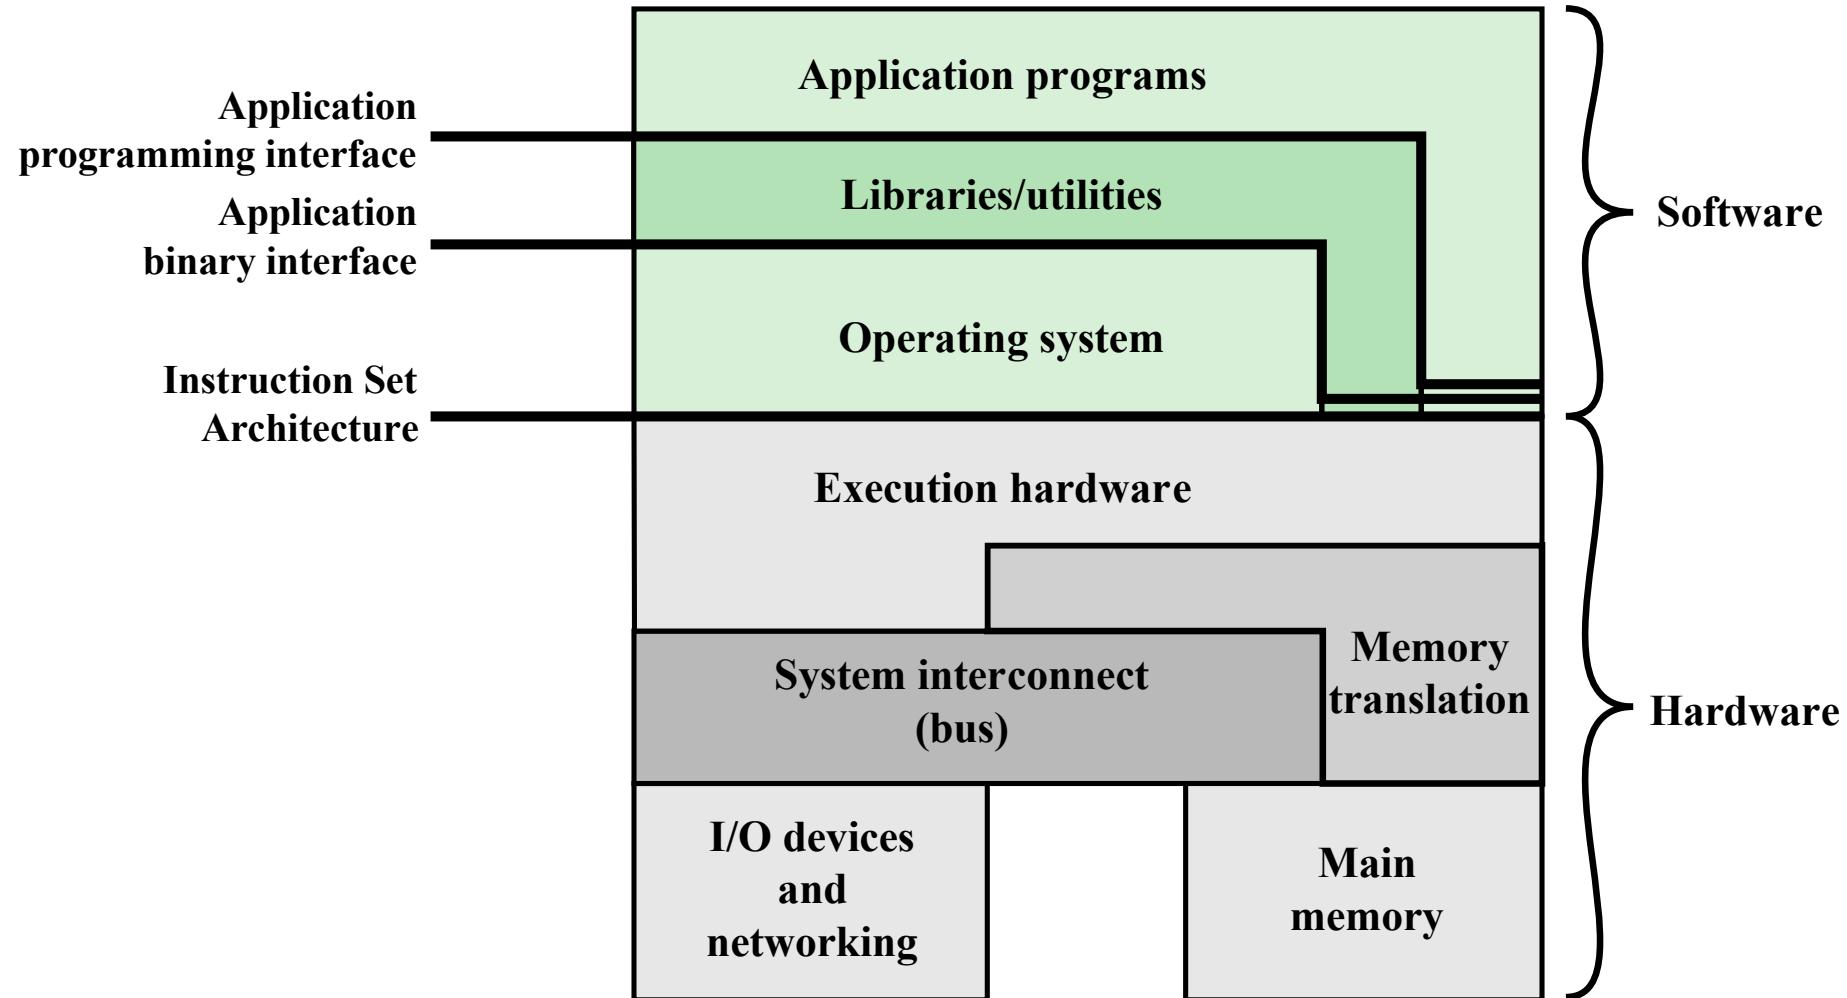


Figure 2.1 Computer Hardware and Software Structure

Operating System Services

- Program development
- Program execution
- Access I/O devices
- Controlled access to files
- System access
- Error detection and response
- Accounting

Key Interfaces

- Instruction set architecture (ISA)
- Application binary interface (ABI)
- Application programming interface (API)

The Role of an OS

- A computer is a set of resources for the movement, storage, and processing of data
- The OS is responsible for managing these resources



Operating System as Software



- Functions in the same way as ordinary computer software
- Program, or suite of programs, executed by the processor
- Frequently relinquishes control and must depend on the processor to allow it to regain control

Computer System

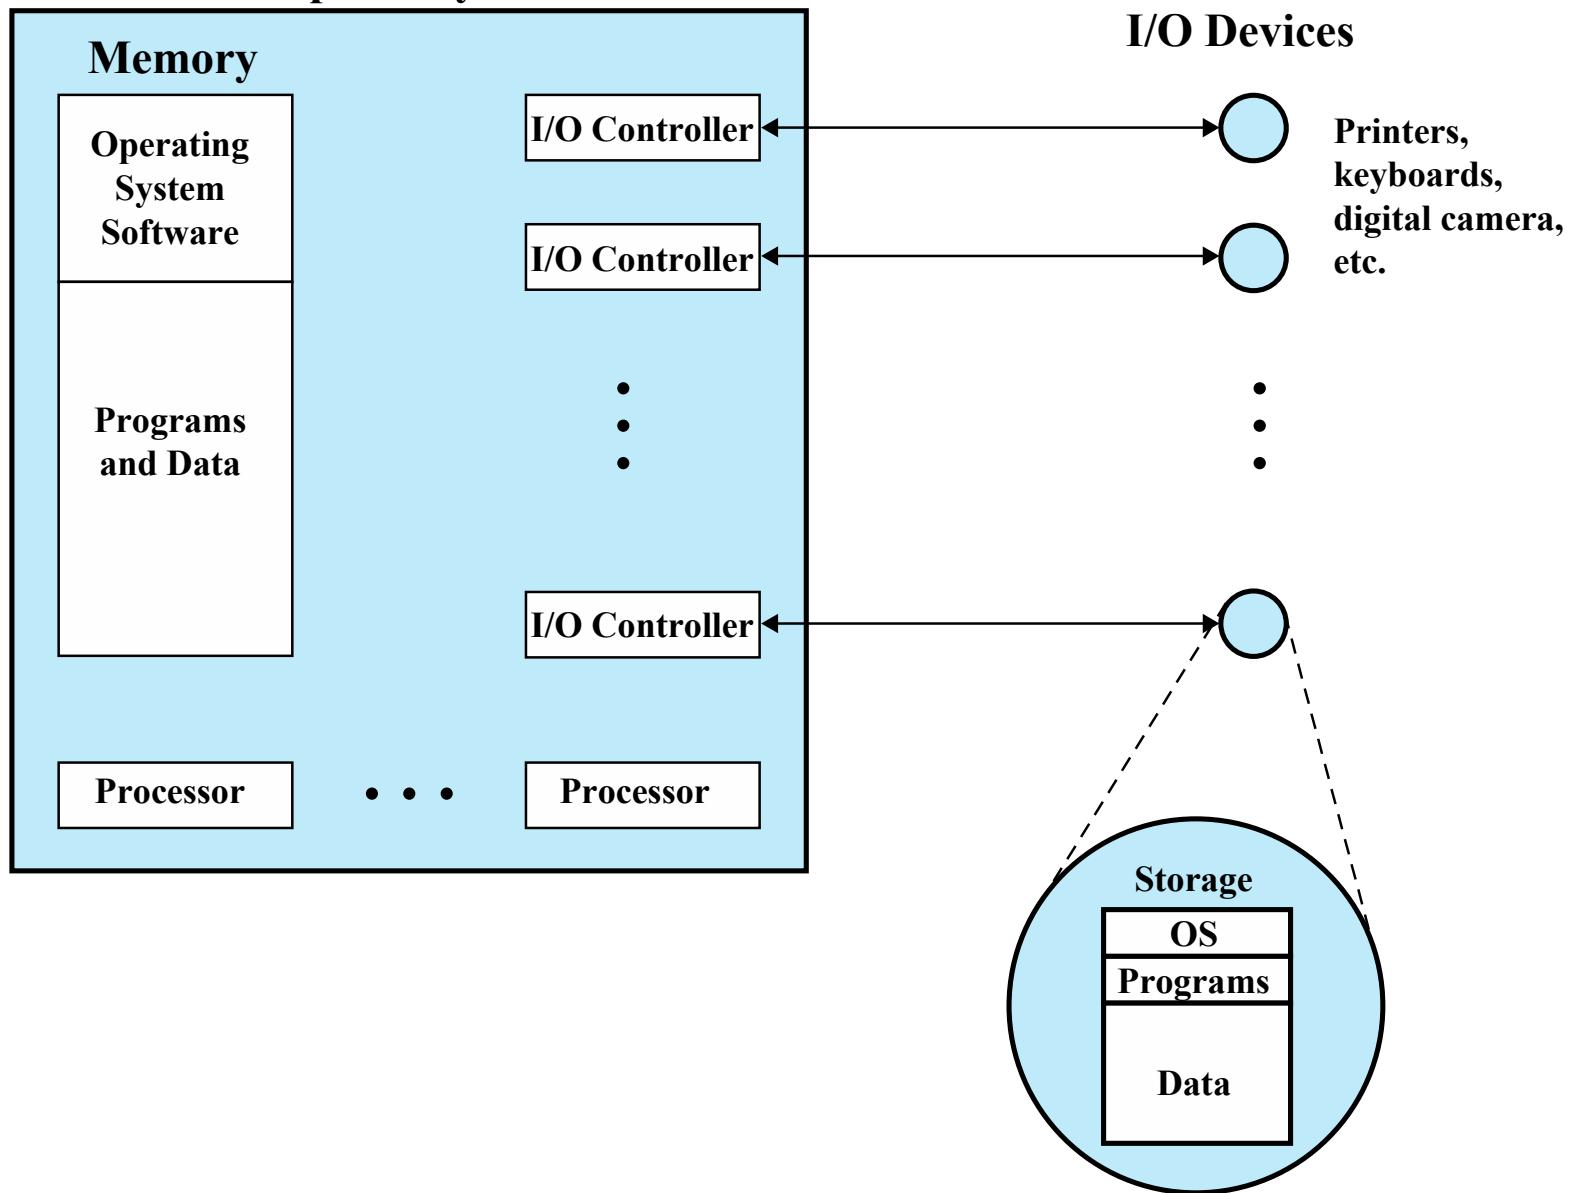


Figure 2.2 The Operating System as Resource Manager

Evolution of Operating Systems

- A major OS will evolve over time for a number of reasons:

hardware upgrades

new types of hardware

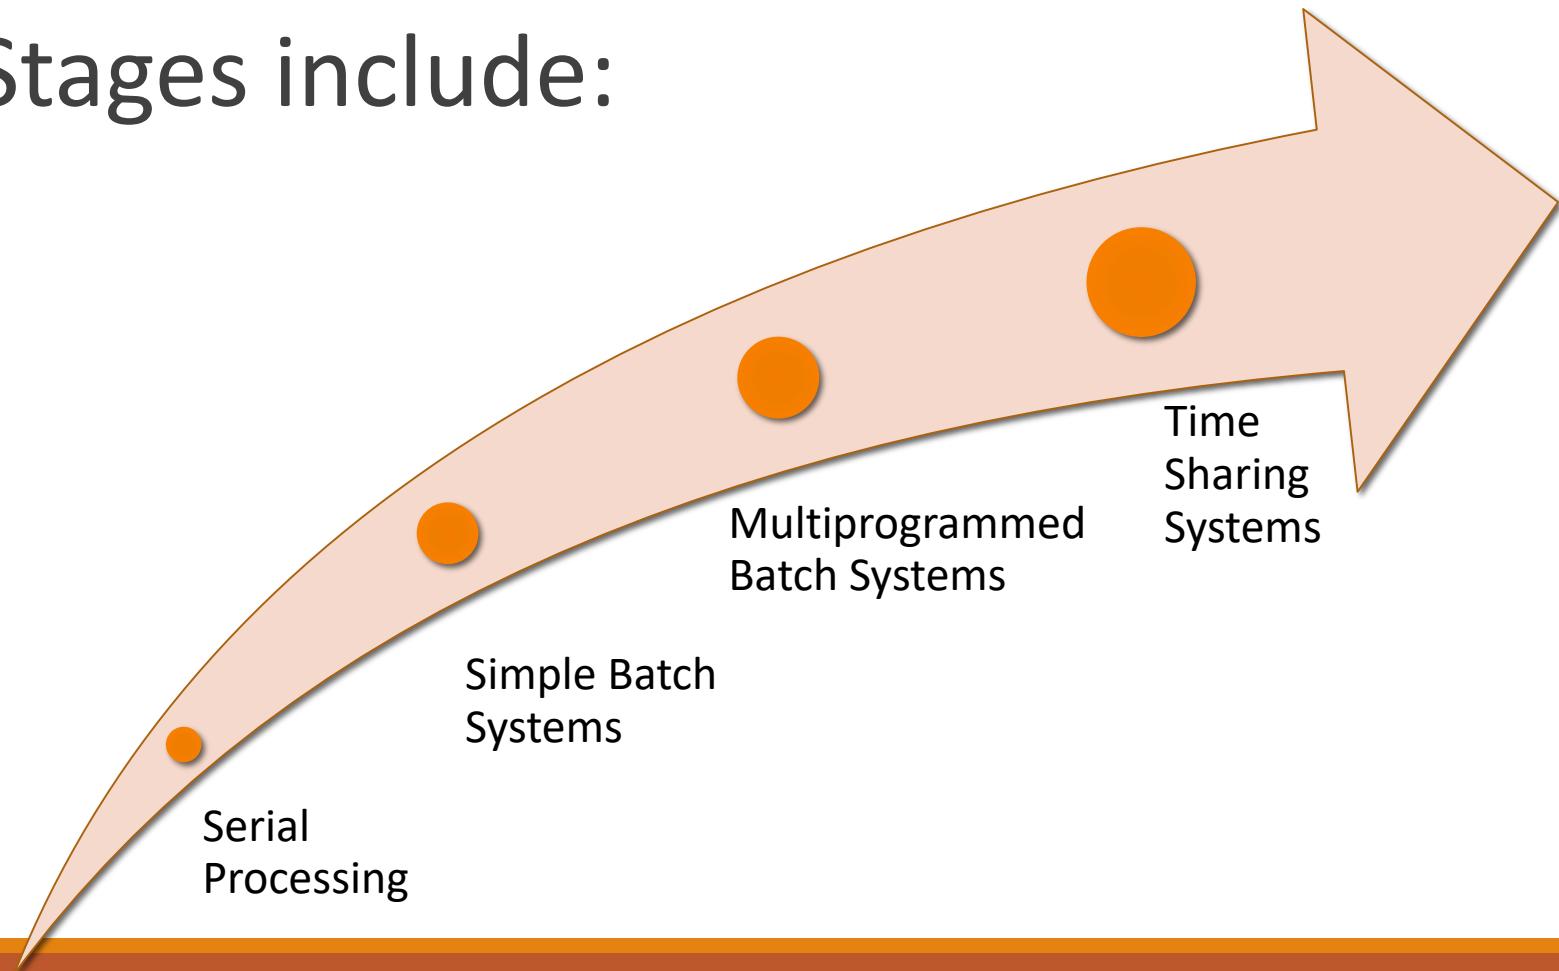
new services

Fixes



Evolution of Operating Systems

- Stages include:



Serial Processing

EARLIEST COMPUTERS:

No operating system

- programmers interacted directly with the computer hardware

Computers ran from a console with display lights, toggle switches, some form of input device, and a printer

Users have access to the computer in “series”

PROBLEMS:

Scheduling:

- most installations used a hardcopy sign-up sheet to reserve computer time
- time allocations could run short or long, resulting in wasted computer time
- Setup time
- a considerable amount of time was spent just on setting up the program to run

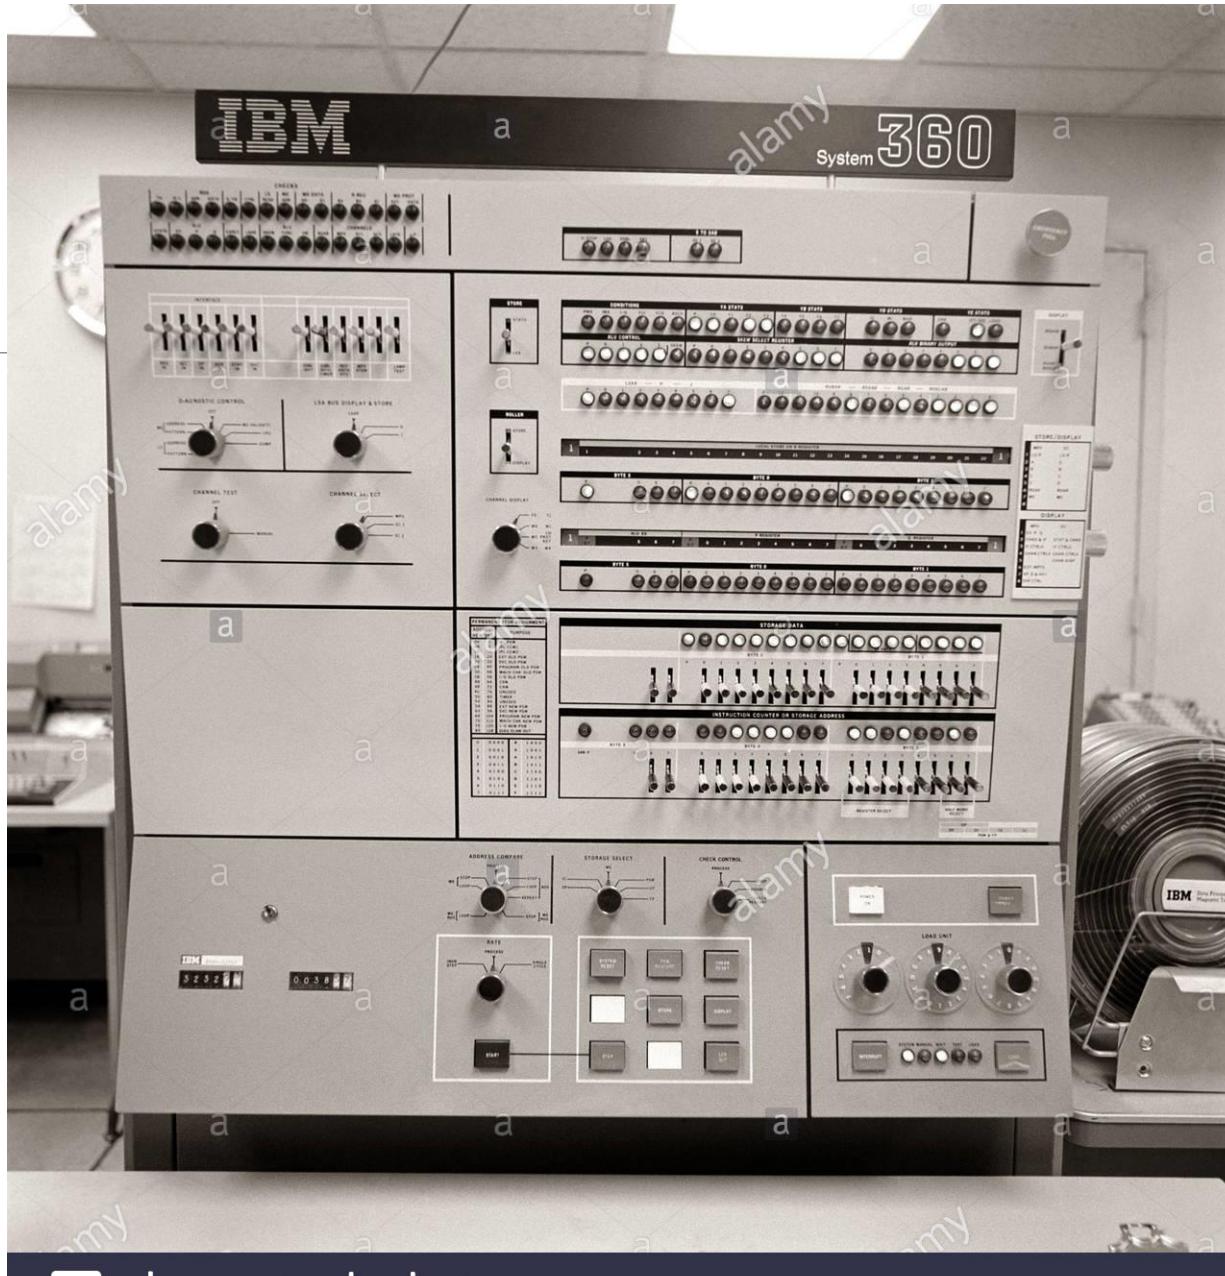
1950s

The computer age dawns:
A Remington Rand "calculating
machine —the first commercially
successful computer using punch
card technology—was born in the
1950s, and was used by the U.S.
Census Bureau.

Credit: Bettmann/Corbis

Simple Batch Systems

- Early computers were very expensive
 - important to maximize processor utilization
- Monitor
 - user no longer has direct access to processor
 - job is submitted to computer operator who batches them together and places them on an input device
 - program branches back to the monitor when finished



a alamy stock photo

AAME8W
www.alamy.com

Monitor Point of View

- Monitor controls the sequence of events
- *Resident Monitor* is software always in memory
- Monitor reads in job and gives control
- Job returns control to monitor

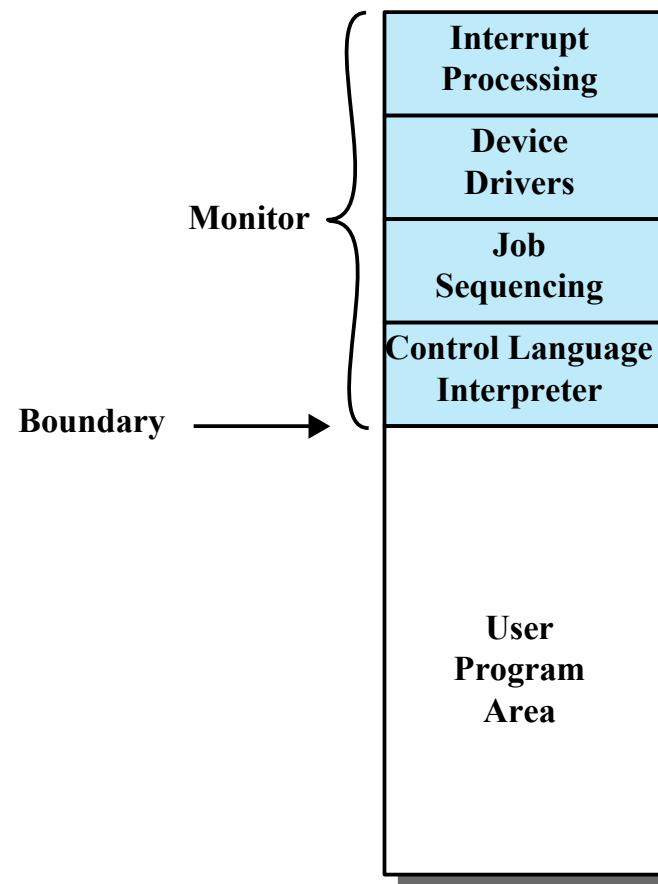


Figure 2.3 Memory Layout for a Resident Monitor

Processor Point of View

- Processor executes instruction from the memory containing the monitor
- Executes the instructions in the user program until it encounters an ending or error condition
- “*control is passed to a job*” means processor is fetching and executing instructions in a user program
- “*control is returned to the monitor*” means that the processor is fetching and executing instructions from the monitor program

Job Control Language (JCL)

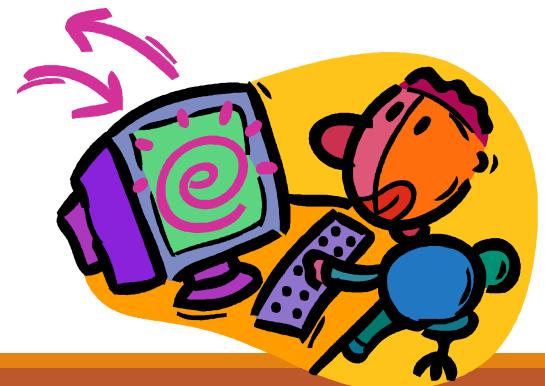
Special type of programming language used to provide instructions to the monitor



what compiler to use



what data to use



Desirable Hardware Features

Memory protection for monitor

- while the user program is executing, it must not alter the memory area containing the monitor

Timer

- prevents a job from monopolizing the system

Privileged instructions

- can only be executed by the monitor

Interrupts

- gives OS more flexibility in controlling user programs

Modes of Operation

User Mode

- user program executes in user mode
- certain areas of memory are protected from user access
- certain instructions may not be executed

Kernel Mode

- monitor executes in kernel mode
- privileged instructions may be executed
- protected areas of memory may be accessed

Simple Batch System Overhead

- Processor time alternates between execution of user programs and execution of the monitor
- Sacrifices:
 - some main memory is now given over to the monitor
 - some processor time is consumed by the monitor
- Despite overhead, the simple batch system improves utilization of the computer

Multiprogrammed Batch Systems

Read one record from file	15 μ s
Execute 100 instructions	1 μ s
Write one record to file	<u>15 μs</u>
TOTAL	31 μ s

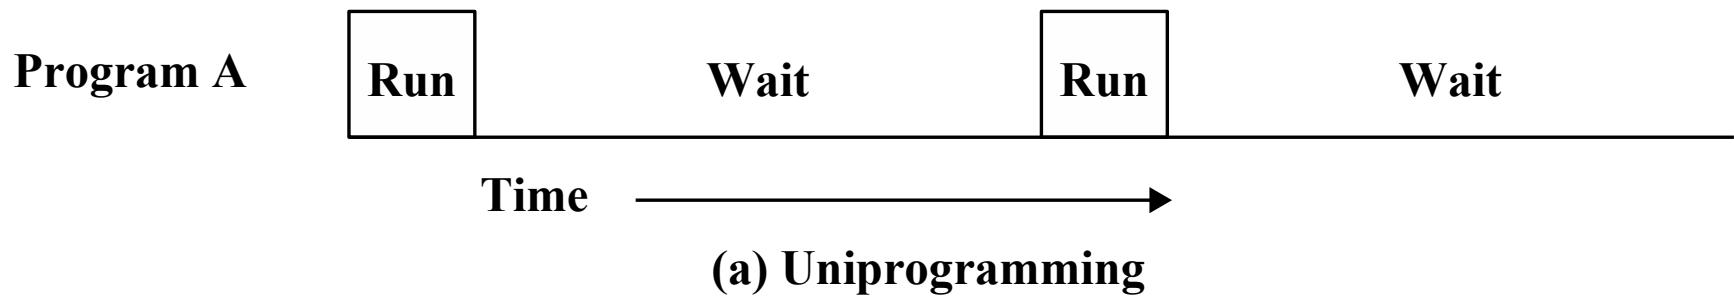
$$\text{Percent CPU Utilization} = \frac{1}{31} = 0.032 = 3.2\%$$

Processor is often idle

- even with automatic job sequencing
- I/O devices are slow compared to processor

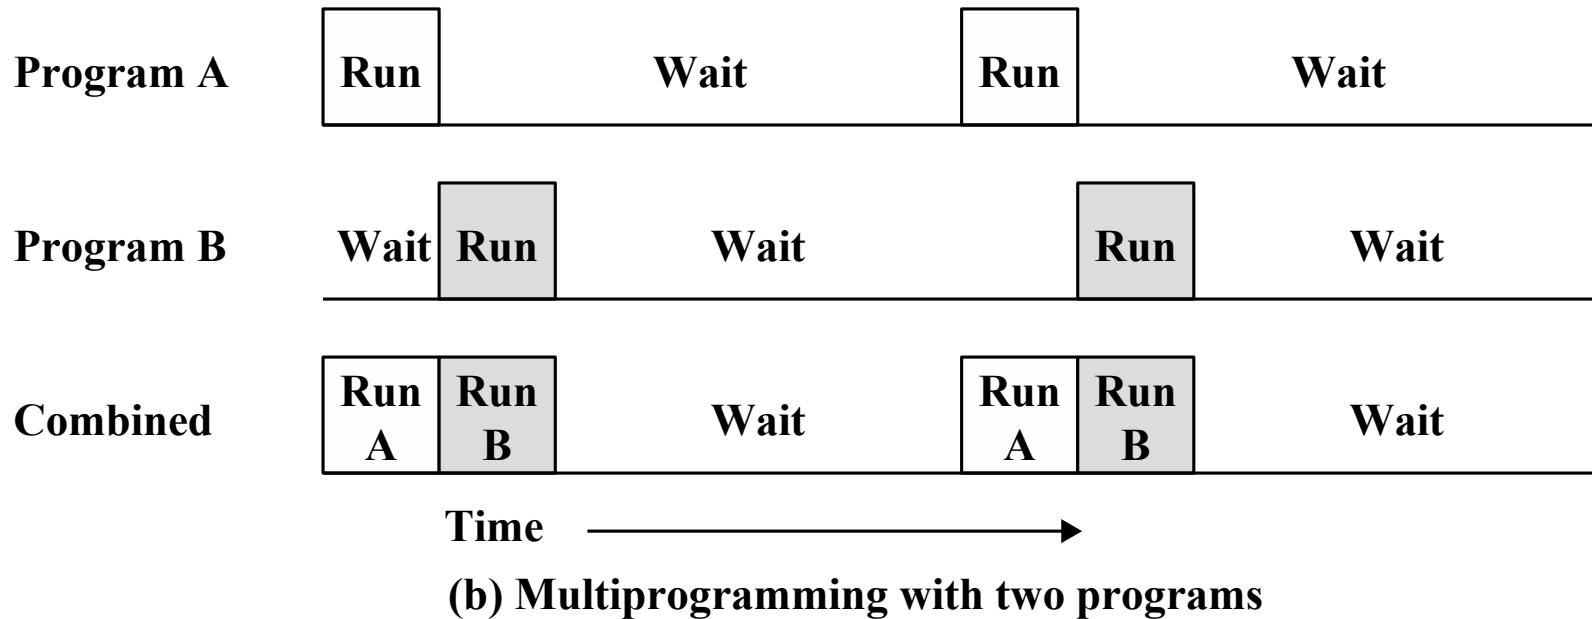
Figure 2.4 System Utilization Example

Uniprogramming



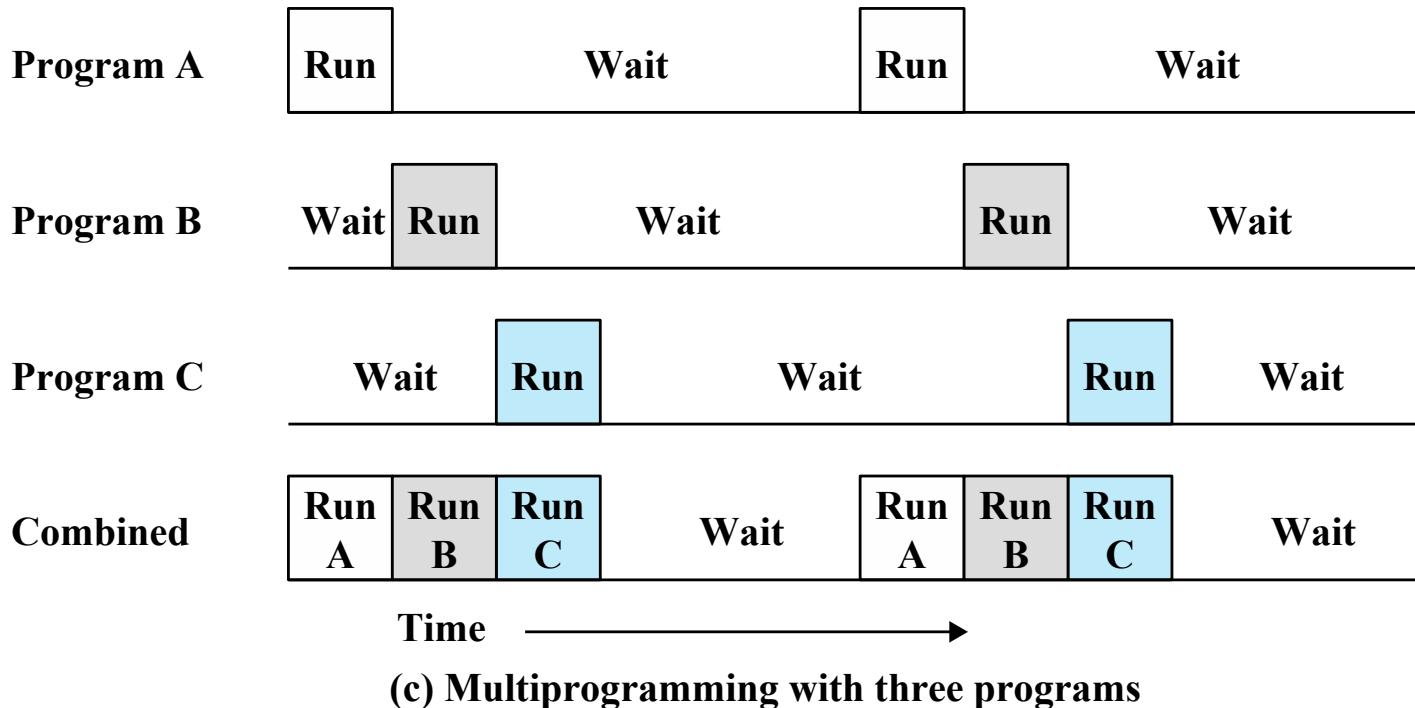
The processor spends a certain amount of time executing, until it reaches an I/O instruction; it must then wait until that I/O instruction concludes before proceeding

Multiprogramming



- There must be enough memory to hold the OS (resident monitor) and one user program
- When one job needs to wait for I/O, the processor can switch to the other job, which is likely not waiting for I/O

Multiprogramming



Multiprogramming

- also known as multitasking
- memory is expanded to hold three, four, or more programs and switch among all of them

Multiprogramming Example

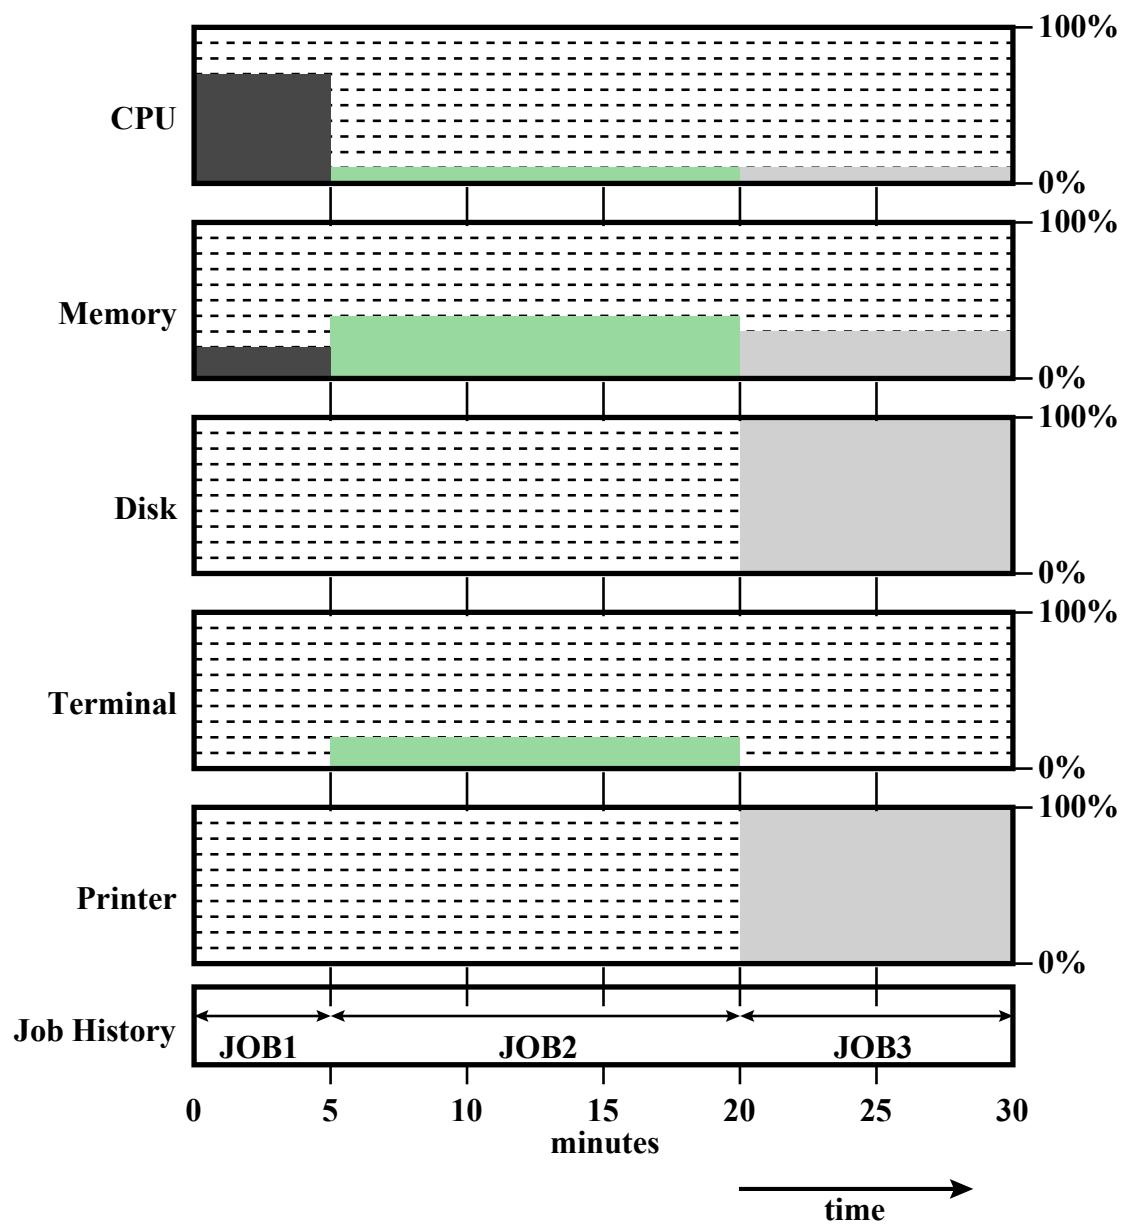
	JOB1	JOB2	JOB3
Type of job	Heavy compute	Heavy I/O	Heavy I/O
Duration	5 min	15 min	10 min
Memory required	50 M	100 M	75 M
Need disk?	No	No	Yes
Need terminal?	No	Yes	No
Need printer?	No	No	Yes

Table 2.1 Sample Program Execution Attributes

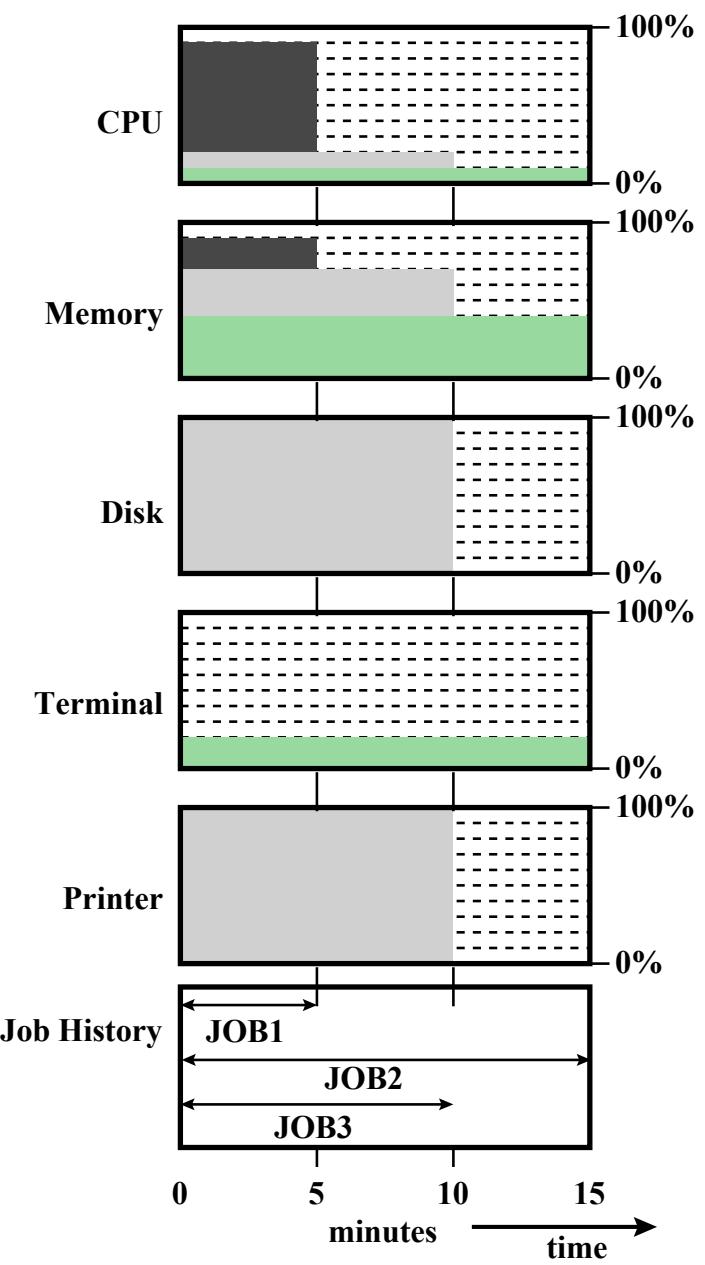
Effects on Resource Utilization

	Uniprogramming	Multiprogramming
Processor use	20%	40%
Memory use	33%	67%
Disk use	33%	67%
Printer use	33%	67%
Elapsed time	30 min	15 min
Throughput	6 jobs/hr	12 jobs/hr
Mean response time	18 min	10 min

Table 2.2 Effects of Multiprogramming on Resource Utilization



(a) Uniprogramming



(b) Multiprogramming

Figure 2.6 Utilization Histograms

Time-Sharing Systems

- Can be used to handle multiple interactive jobs
- Processor time is shared among multiple users
- Multiple users simultaneously access the system through terminals, with the OS interleaving the execution of each user program in a short burst or quantum of computation

Batch Multiprogramming vs. Time Sharing

	Batch Multiprogramming	Time Sharing
Principal objective	Maximize processor use	Minimize response time
Source of directives to operating system	Job control language commands provided with the job	Commands entered at the terminal

Table 2.3 Batch Multiprogramming versus Time Sharing

Compatible Time-Sharing Systems

CTSS

One of the first time-sharing operating systems

- Developed at MIT by a group known as Project MAC
- Ran on a computer with 32,000 36-bit words of main memory, with the resident monitor consuming 5000 of that
- To simplify both the monitor and memory management a program was always loaded to start at the location of the 5000th word

TIME SLICING

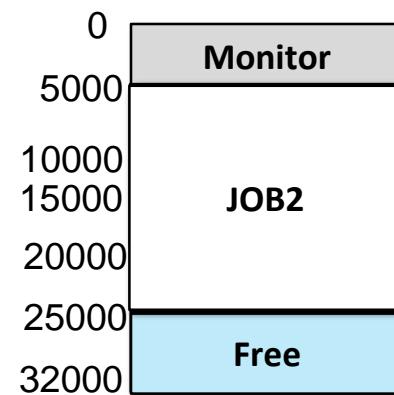
- System clock generates interrupts at a rate of approximately one every 0.2 seconds
- At each interrupt OS regained control and could assign processor to another user
- At regular time intervals the current user would be preempted and another user loaded in
- Old user programs and data were written out to disk
- Old user program code and data were restored in main memory when that program was next given a turn

JOB Sizes:

- JOB1: 15,000
- JOB2: 20,000
- JOB3: 5,000
- JOB4: 10,000

Sequence:

- JOB 1
- JOB 2
- JOB 3
- JOB 1
- JOB 4
- JOB 2



Major Achievements

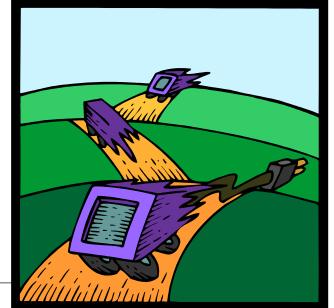
Operating Systems are among the most complex pieces of software ever developed



Major advances in development include:

- processes
- memory management
- information protection and security
- scheduling and resource management
- system structure

Process



Fundamental to the structure of operating systems

A *process* can be defined as:

a program in execution

an instance of a running program

the entity that can be assigned to, and executed on, a processor

a unit of activity characterized by a single sequential thread of execution, a current state, and an associated set of system resources

Development of the Process

- Three major lines of computer system development created problems in timing and synchronization that contributed to the development:

multiprogramming batch operation

- processor is switched among the various programs residing in main memory

time sharing

- be responsive to the individual user but be able to support many users simultaneously

real-time transaction systems

- a number of users are entering queries or updates against a database

Causes of Errors

Improper synchronization

- a program must wait until the data are available in a buffer
- improper design of the signaling mechanism can result in loss or duplication

Failed mutual exclusion

- more than one user or program attempts to make use of a shared resource at the same time
- only one routine at a time allowed to perform an update against the file

Nondeterminate program operation

- program execution is interleaved by the processor when memory is shared
- the order in which programs are scheduled may affect their outcome

Deadlocks

- it is possible for two or more programs to be hung up waiting for each other
- may depend on the chance timing of resource allocation and release

Components of a Process

A process contains three components:

- an executable program
- the associated data needed by the program (variables, work space, buffers, etc.)
- the execution context (or “process state”) of the program

- The execution context is essential:
 - it is the internal data by which the OS is able to supervise and control the process
 - includes the contents of the various process registers
 - includes information such as the priority of the process and whether the process is waiting for the completion of a particular I/O event

Process Management

- The entire state of the process at any instant is contained in its context
- New features can be designed and incorporated into the OS by expanding the context to include any new information needed to support the feature

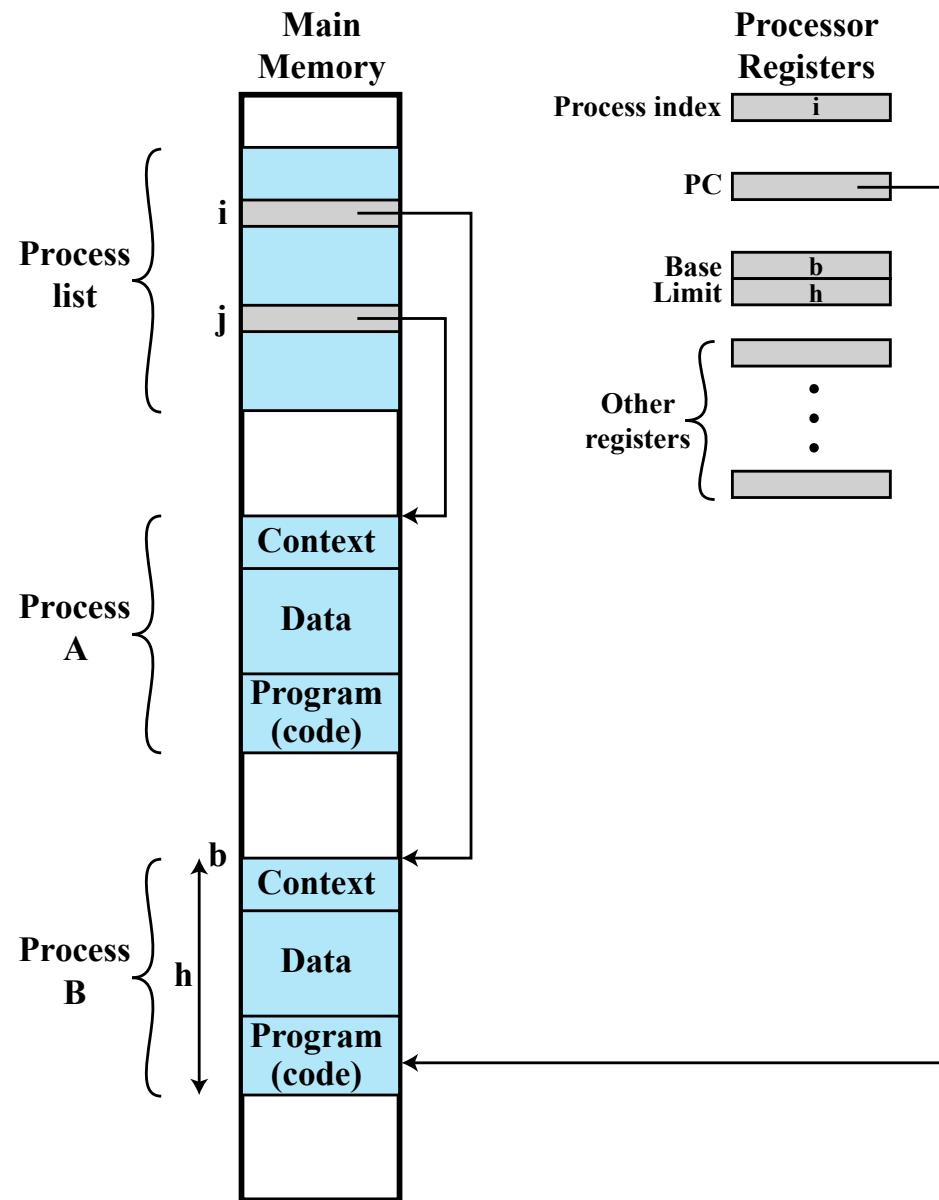


Figure 2.8 Typical Process Implementation

Memory Management

The OS has **five** principal storage management responsibilities:

process isolation

automatic allocation
and management

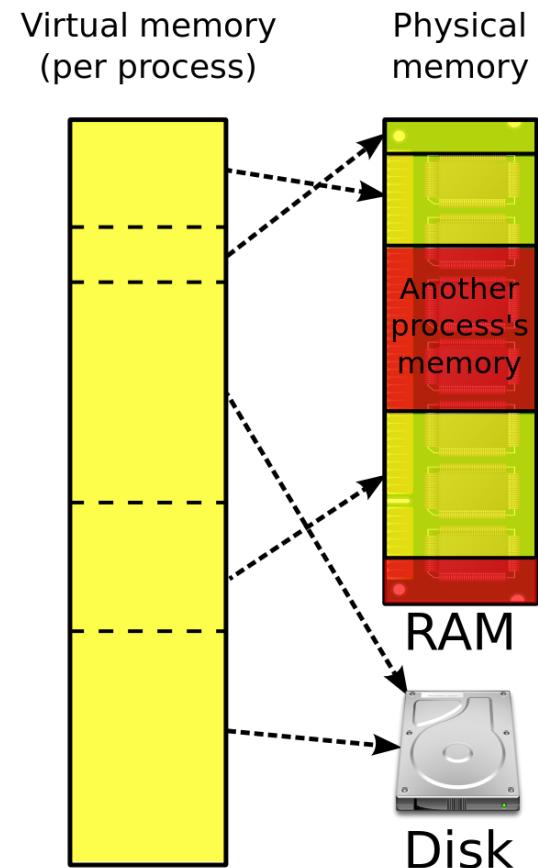
support of modular
programming

protection and access control

long-term storage

Virtual Memory

- A facility that allows programs to address memory from a logical point of view, without regard to the amount of main memory physically available
- Conceived to meet the requirement of having multiple user jobs reside in main memory concurrently



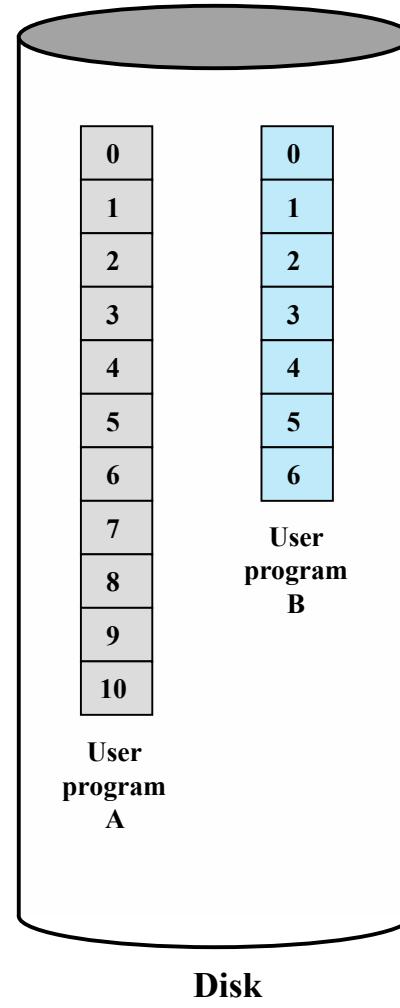
Paging

- Allows processes to be comprised of a number of fixed-size blocks, called pages
- Program references a word by means of a virtual address
 - consists of a page number and an offset within the page
 - each page may be located anywhere in main memory
- Provides for a dynamic mapping between the virtual address used in the program and a real (or physical) address in main memory

A.1			
	A.0	A.2	
	A.5		
B.0	B.1	B.2	B.3
		A.7	
	A.9		
		A.8	
B.5	B.6		

Main Memory

Main memory consists of a number of fixed-length frames, each equal to the size of a page. For a program to execute, some or all of its pages must be in main memory.



Disk

Secondary memory (disk) can hold many fixed-length pages. A user program consists of some number of pages. Pages for all programs plus the operating system are on disk, as are files.

Figure 2.9 Virtual Memory Concepts

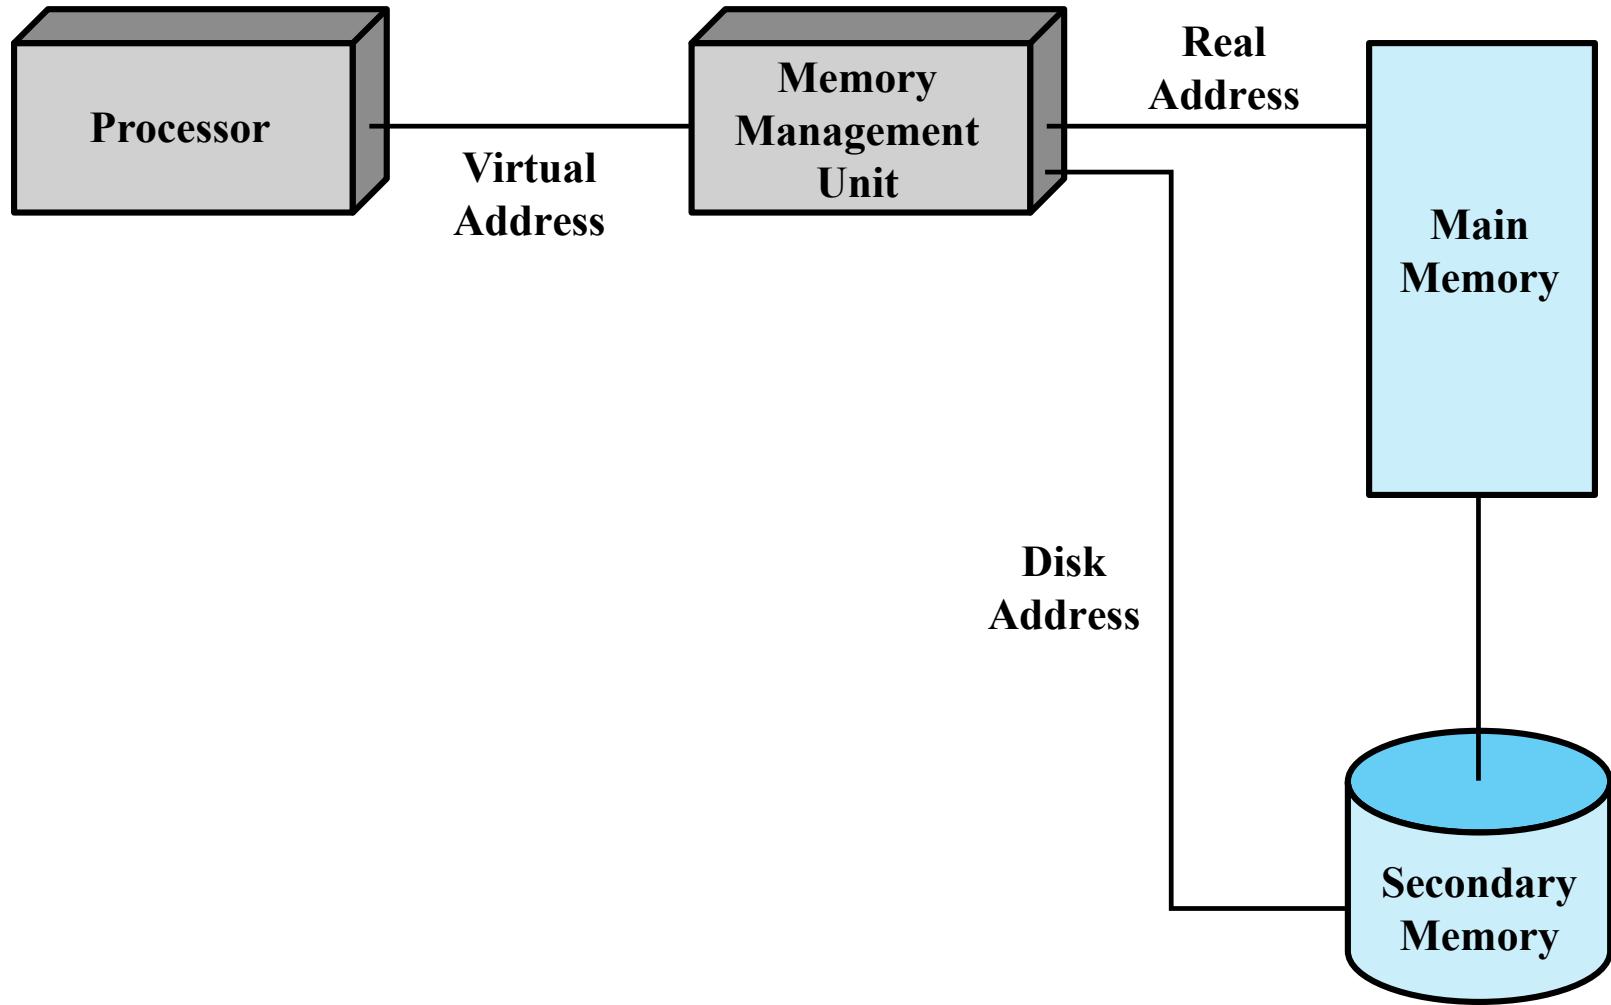


Figure 2.10 Virtual Memory Addressing

Information Protection and Security

- The nature of the threat that concerns an organization will vary greatly depending on the circumstances
- The problem involves controlling access to computer systems and the information stored in them



Scheduling and Resource Management

- Key responsibility of an OS is managing resources
- Resource allocation policies must consider:



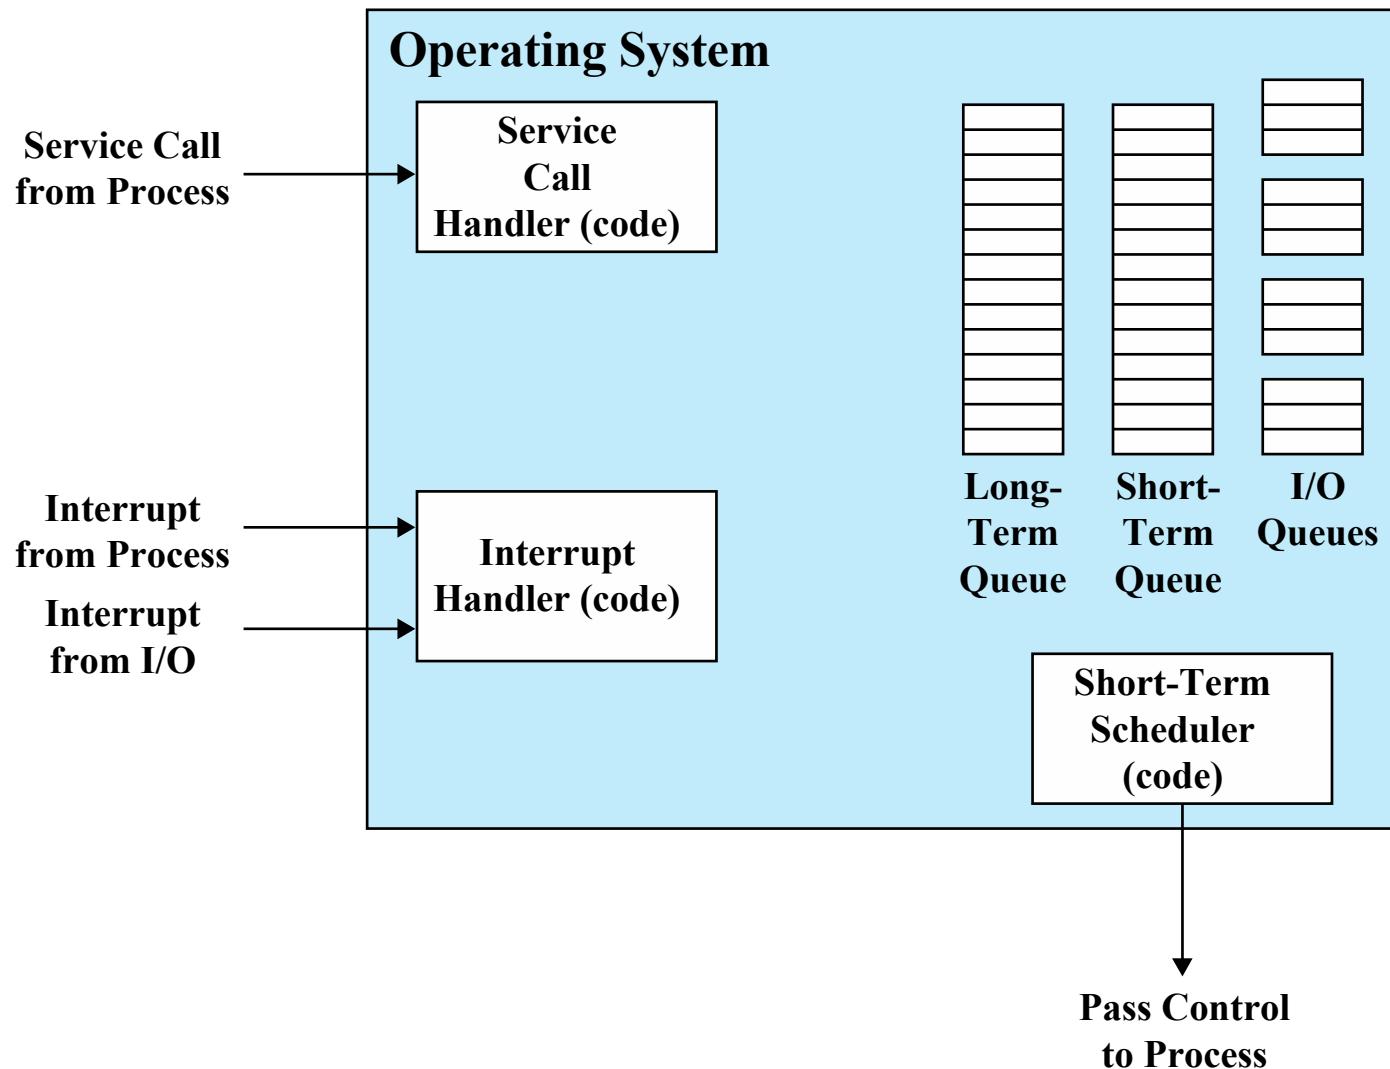


Figure 2.11 Key Elements of an Operating System for Multiprogramming

Different Architectural Approaches

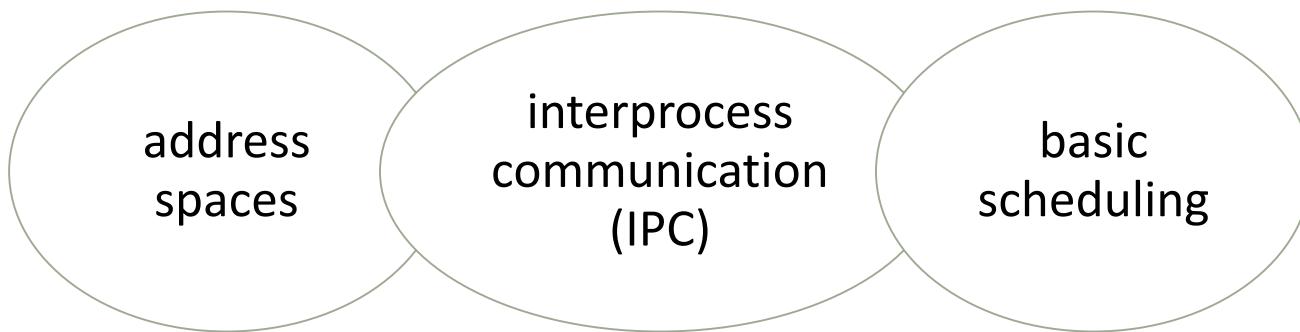
Demands on operating systems require new ways of organizing the OS

Different approaches and design elements have been tried:

- microkernel architecture
- multithreading
- symmetric multiprocessing
- distributed operating systems
- object-oriented design

Microkernel Architecture

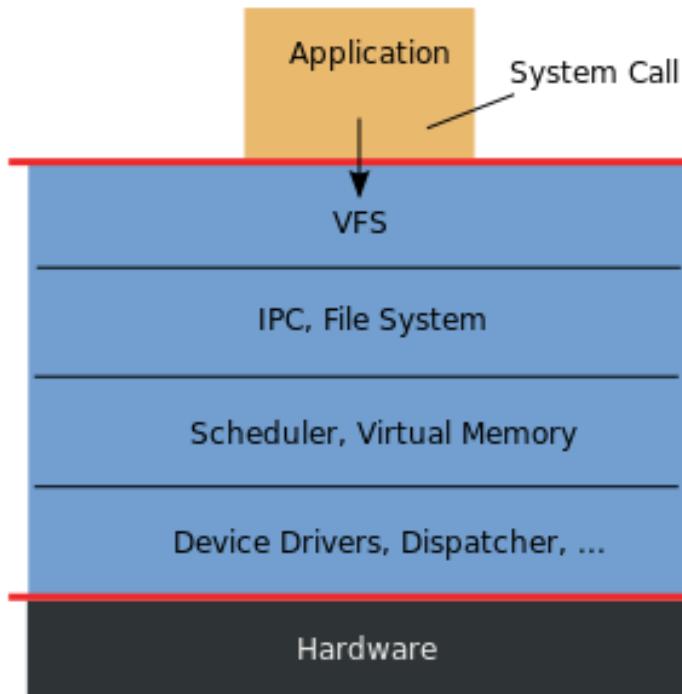
Assigns only a few essential functions to the kernel:



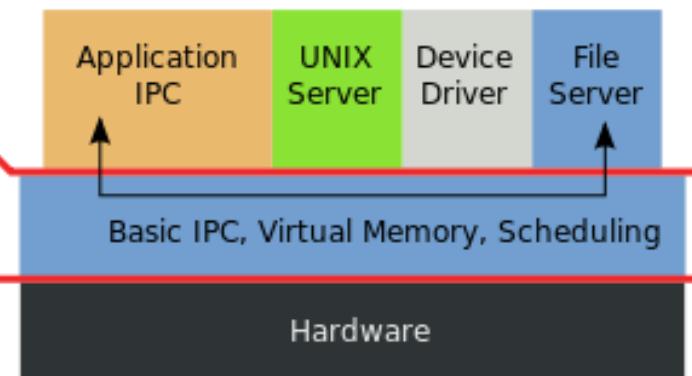
- Other OS services are provided by processes run in user mode called servers.
- Servers are customized to specific application or environment.

Microkernel Architecture

Monolithic Kernel
based Operating System



Microkernel
based Operating System



Microkernel Architecture

- The approach:



simplifies implementation

provides flexibility

is well suited to
a distributed
environment

Multithreading

Technique in which a process, executing an application, is divided into threads that can run concurrently

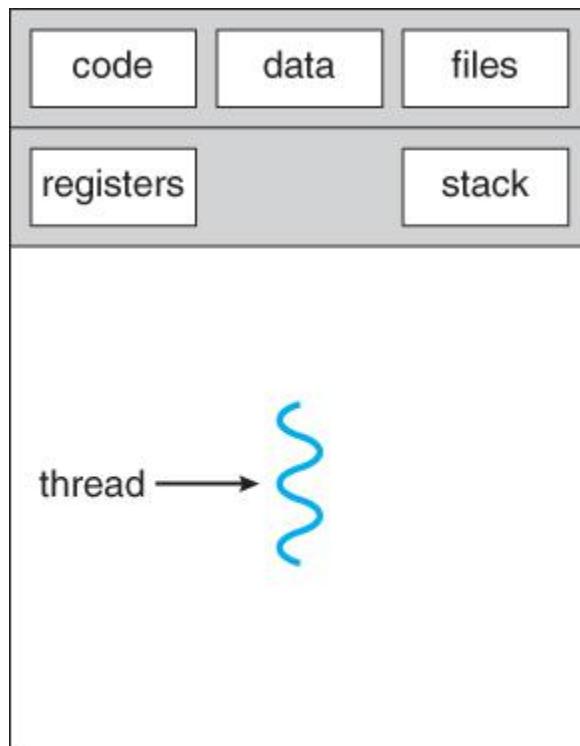
Thread

- dispatchable unit of work
- includes a processor context and its own data area to enable subroutine branching
- executes sequentially and is interruptible

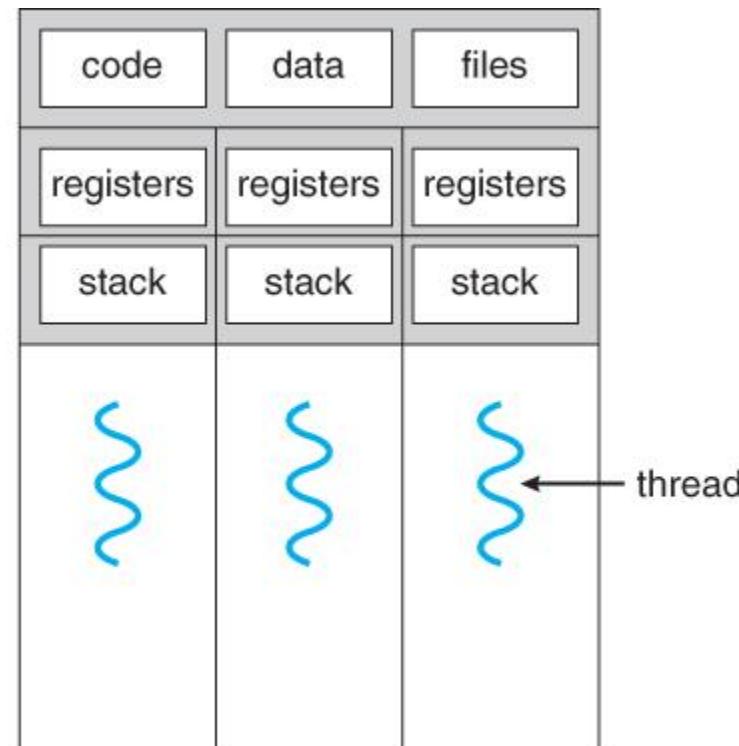
Process

- a collection of one or more threads and associated system resources
- programmer has greater control over the modularity of the application and the timing of application related events

Multithreading



single-threaded process



multithreaded process

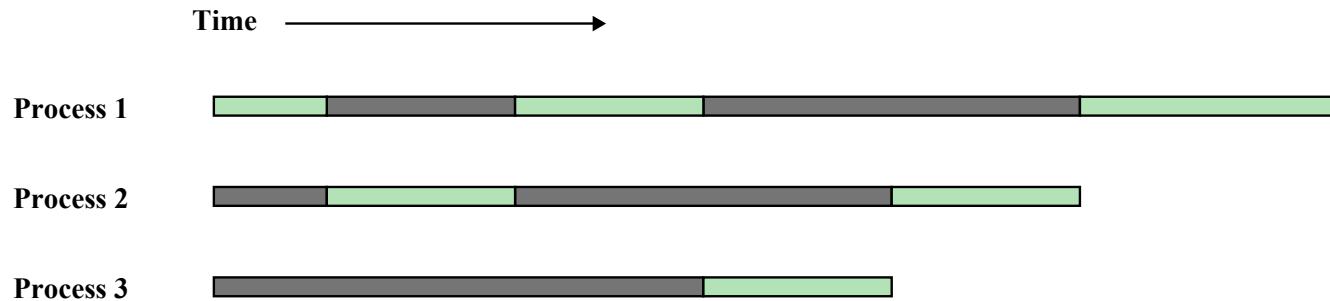
Symmetric Multiprocessing (SMP)

- Term that refers to a computer hardware architecture and also to the OS behavior that exploits that architecture
- Several processes can run in parallel
- Multiple processors are transparent to the user
 - these processors share same main memory and I/O facilities
 - all processors can perform the same functions
- The OS takes care of scheduling of threads or processes on individual processors and of synchronization among processors

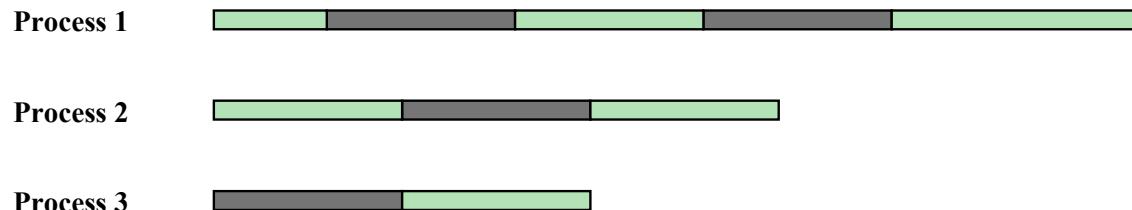
SMP Advantages

- Performance** more than one process can be running simultaneously, each on a different processor
- Availability** failure of a single process does not halt the system
- Incremental Growth** performance of a system can be enhanced by adding an additional processor
- Scaling** vendors can offer a range of products based on the number of processors configured in the system





(a) Interleaving (multiprogramming, one processor)



(b) Interleaving and overlapping (multiprocessing; two processors)

 Blocked  Running

Figure 2.12 Multiprogramming and Multiprocessing

OS Design

DISTRIBUTED OPERATING SYSTEM

- Provides the illusion of
 - a single main memory space
 - single secondary memory space
 - unified access facilities
- State of the art for distributed operating systems lags that of uniprocessor and SMP operating systems

OBJECT-ORIENTED DESIGN

- Used for adding modular extensions to a small kernel
- Enables programmers to customize an operating system without disrupting system integrity
- Eases the development of distributed tools and full-blown distributed operating systems

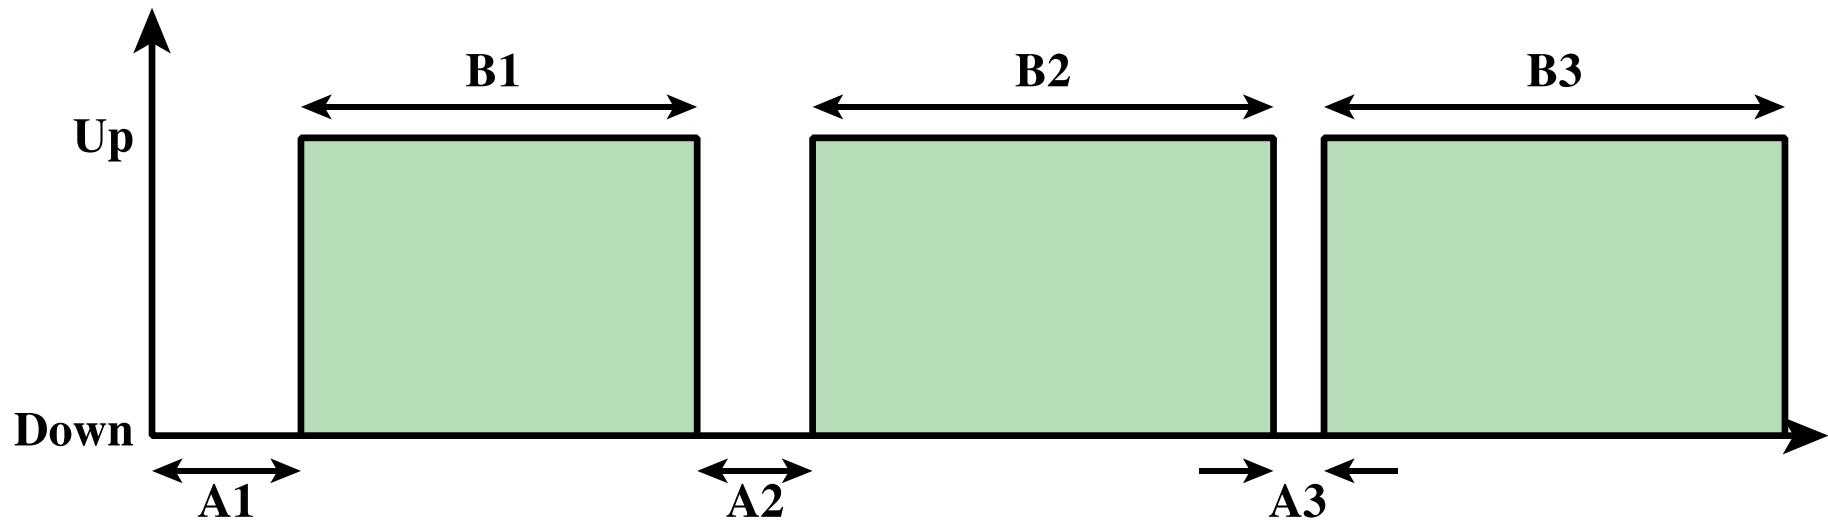
Fault Tolerance

- Refers to the ability of a system or component to continue normal operation despite the presence of hardware or software faults
- Typically involves some degree of redundancy
- Intended to increase the reliability of a system
 - typically comes with a cost in financial terms or performance
- The extent adoption of fault tolerance measures must be determined by how critical the resource is

Fundamental Concepts

The basic measures are:

- Reliability
 - $R(t)$
 - defined as the probability of its correct operation up to time t given that the system was operating correctly at time $t=0$
- Mean time to failure (MTTF)
 - mean time to repair (MTTR) is the average time it takes to repair or replace a faulty element
- Availability
 - defined as the fraction of time the system is available to service users' requests



$$MTTF = \frac{B1 + B2 + B3}{3}$$

$$MTTR = \frac{A1 + A2 + A3}{3}$$

$$Avail = \frac{MTTF}{MTTF + MTTR}$$

Figure 2.13 System Operational States

Availability Classes

Class	Availability	Annual Downtime
Continuous	1.0	0
Fault Tolerant	0.99999	5 minutes
Fault Resilient	0.9999	53 minutes
High Availability	0.999	8.3 hours
Normal Availability	0.99 - 0.995	44-87 hours

Table 2.4 Availability Classes

Fault Categories

Permanent

- a fault that, after it occurs, is always present
- the fault persists until the faulty component is replaced or repaired

Temporary

- a fault that is not present all the time for all operating conditions
- can be classified as
 - Transient – a fault that occurs only once
 - Intermittent – a fault that occurs at multiple, unpredictable times

Spatial (physical) redundancy

involves the use of multiple components that either perform the same function simultaneously or are configured so that one component is available as a backup in case of the failure of another component

Temporal redundancy

involves repeating a function or operation when an error is detected
effective with temporary faults but not useful for permanent faults

Information redundancy

provides fault tolerance by replicating or coding data in such a way that bit errors can be both detected and corrected

Operating System Mechanisms

A number of techniques can be incorporated into OS software to support fault tolerance:

- process isolation
- concurrency
- virtual machines
- checkpoints and rollbacks

Symmetric Multiprocessor OS Considerations

A multiprocessor OS must provide all the functionality of a multiprogramming system plus additional features to accommodate multiple processors

Key design issues:

Simultaneous concurrent processes or threads

kernel routines need to be reentrant to allow several processors to execute the same kernel code simultaneously

Scheduling

any processor may perform scheduling, which complicates the task of enforcing a scheduling policy

Synchronization

with multiple active processes having potential access to shared address spaces or shared I/O resources, care must be taken to provide effective synchronization

Memory management

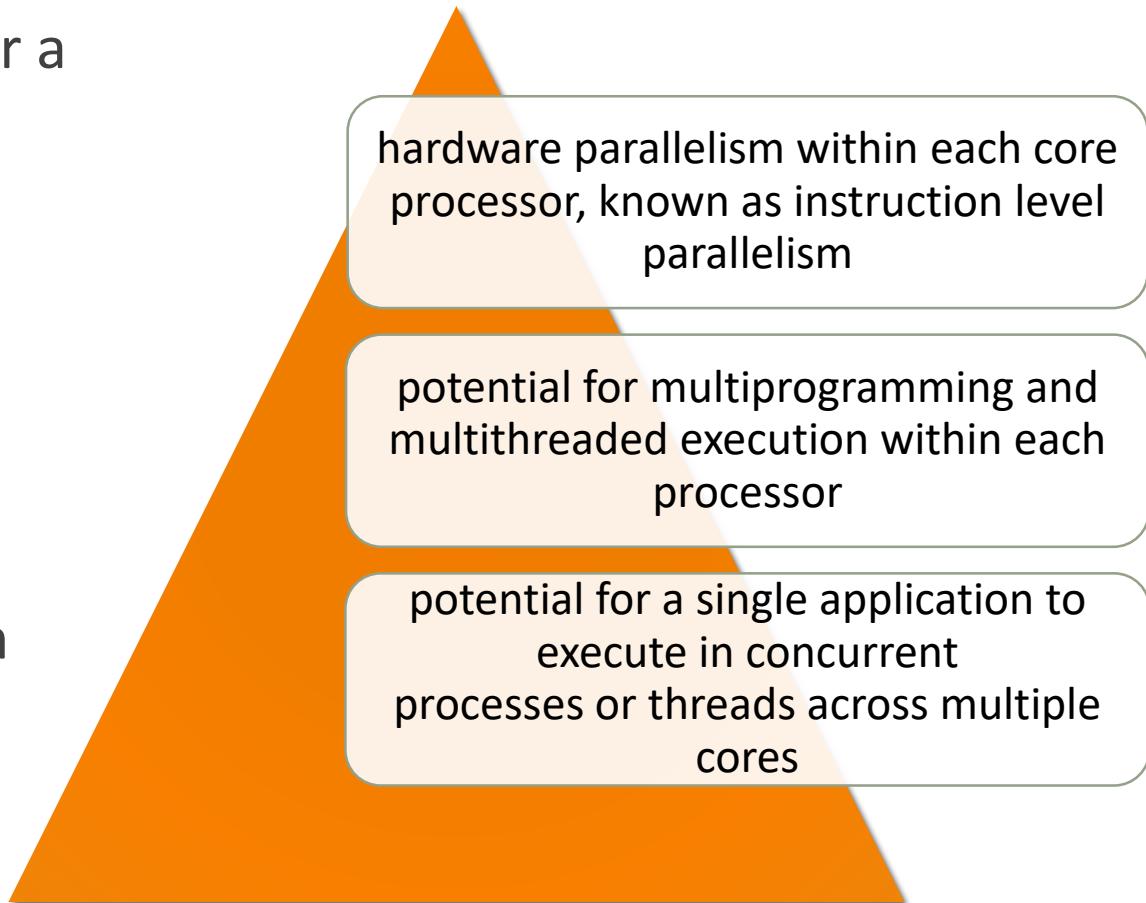
the reuse of physical pages is the biggest problem of concern

Reliability and fault tolerance

the OS should provide graceful degradation in the face of processor failure

Multicore OS Considerations

- The design challenge for a many-core multicore system is to efficiently harness the multicore processing power and intelligently manage the substantial on-chip resources efficiently
- Potential for parallelism exists at three levels:



Grand Central Dispatch

Developer must decide what pieces can or should be executed simultaneously or in parallel

Grand Central Dispatch (GCD)

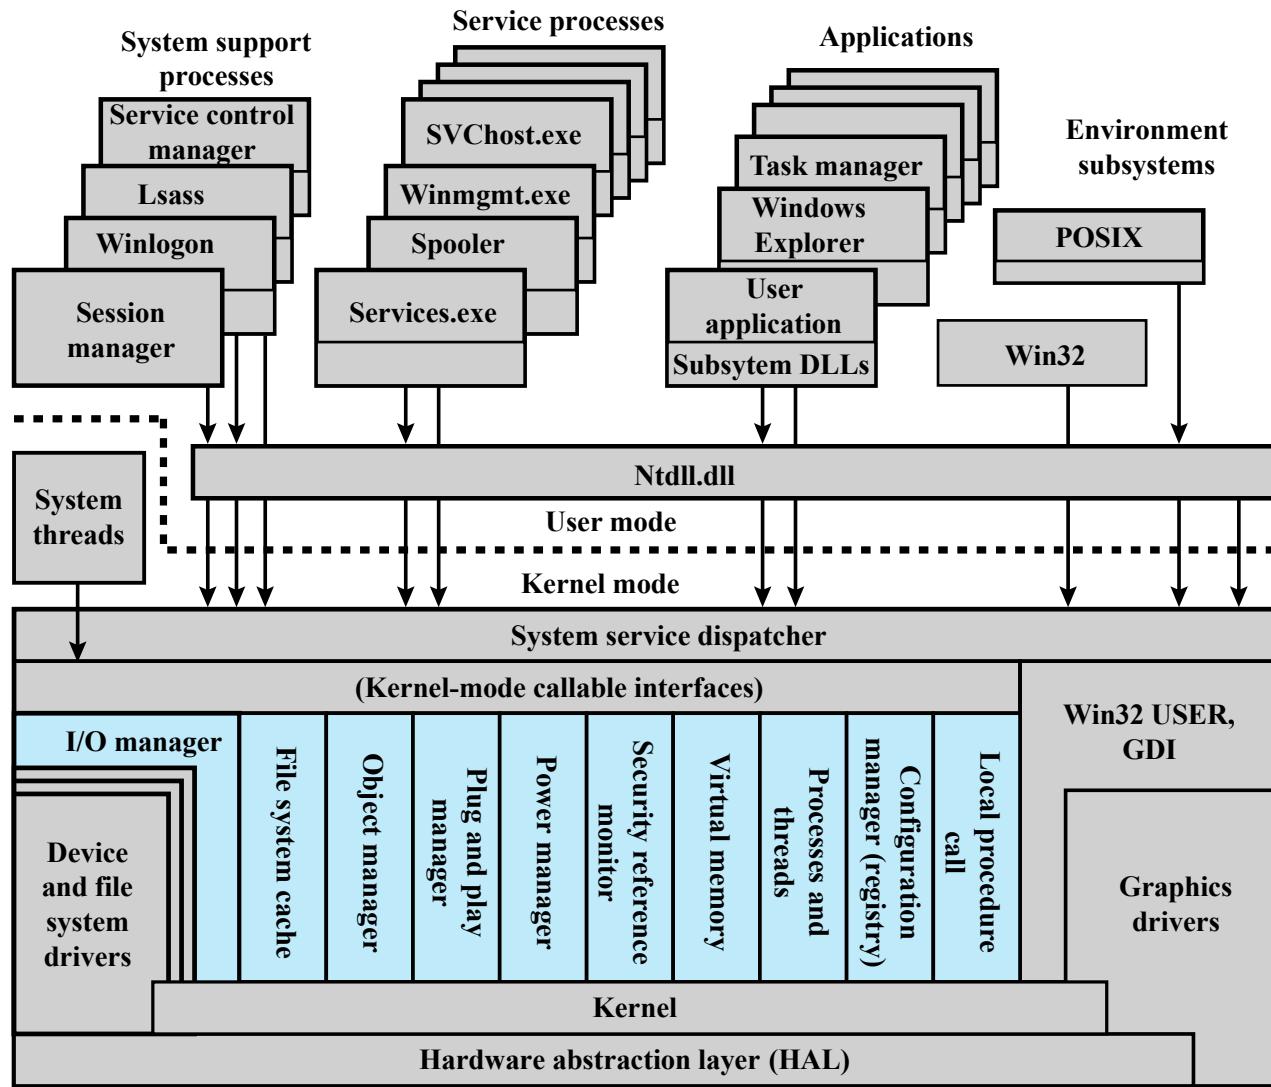
- implemented in Mac Os X 10.6
- helps a developer once something has been identified that can be split off into a separate task
- thread pool mechanism
- allows anonymous functions as a way of specifying tasks

Virtual Machine Approach

Allows one or more cores to be dedicated to a particular process and then leave the processor alone to devote its efforts to that process

Multicore OS could then act as a hypervisor that makes a high-level decision to allocate cores to applications but does little in the way of resource allocation beyond that





Lsass = local security authentication server
 POSIX = portable operating system interface
 GDI = graphics device interface
 DLL = dynamic link libraries

Colored area indicates Executive

Figure 2.14 Windows Architecture

Kernel-Mode Components of Windows

Executive

- contains the core OS services

Kernel

- controls execution of the processors

Hardware Abstraction Layer (HAL)

- maps between generic hardware commands and responses and those unique to a specific platform

Device Drivers

- dynamic libraries that extend the functionality of the Executive

Windowing and Graphics System

- implements the GUI functions

User-Mode Processes

Four basic types are supported by Windows:

Special System Processes

- user-mode services needed to manage the system

Service Processes

- the printer spooler, event logger, and user-mode components that cooperate with device drivers, and various network services

Environment Subsystems

- provide different OS personalities (environments)

User Applications

- executables (EXEs) and DLLs that provide the functionality users run to make use of the system

Client/Server Model

Windows OS services, environmental subsystems, and applications are all structured using the client/server model

- Common in distributed systems, but can be used internal to a single system
- Processes communicate via RPC

Advantages:

- it simplifies the Executive
- it improves reliability
- it provides a uniform means for applications to communicate with services via RPCs without restricting flexibility
- it provides a suitable base for distributed computing

Threads and SMP

Two important characteristics of Windows are its support for threads and for symmetric multiprocessing (SMP)

- OS routines can run on any available processor, and different routines can execute simultaneously on different processors
- Windows supports the use of multiple threads of execution within a single process. Multiple threads within the same process may execute on different processors simultaneously
- server processes may use multiple threads to process requests from more than one client simultaneously
- Windows provides mechanisms for sharing data and resources between processes and flexible interprocess communication capabilities

Windows Objects

Windows draws heavily on the concepts of object-oriented design

Key object-oriented concepts used by Windows are:



Asynchronous Procedure Call	Used to break into the execution of a specified thread and to cause a procedure to be called in a specified processor mode.
Deferred Procedure Call	Used to postpone interrupt processing to avoid delaying hardware interrupts. Also used to implement timers and inter-processor communication
Interrupt	Used to connect an interrupt source to an interrupt service routine by means of an entry in an Interrupt Dispatch Table (IDT). Each processor has an IDT that is used to dispatch interrupts that occur on that processor.
Process	Represents the virtual address space and control information necessary for the execution of a set of thread objects. A process contains a pointer to an address map, a list of ready threads containing thread objects, a list of threads belonging to the process, the total accumulated time for all threads executing within the process, and a base priority.
Thread	Represents thread objects, including scheduling priority and quantum, and which processors the thread may run on.
Profile	Used to measure the distribution of run time within a block of code. Both user and system code can be profiled.

Table 2.5 Windows Kernel Control Objects

Traditional UNIX Systems

Were developed at Bell Labs and became operational on a PDP-7 in 1970

Incorporated many ideas from Multics

PDP-11 was a milestone because it first showed that UNIX would be an OS for all computers

Next milestone was rewriting UNIX in the programming language C

- demonstrated the advantages of using a high-level language for system code
- Was described in a technical journal for the first time in 1974
- First widely available version outside Bell Labs was Version 6 in 1976
- Version 7, released in 1978 is the ancestor of most modern UNIX systems

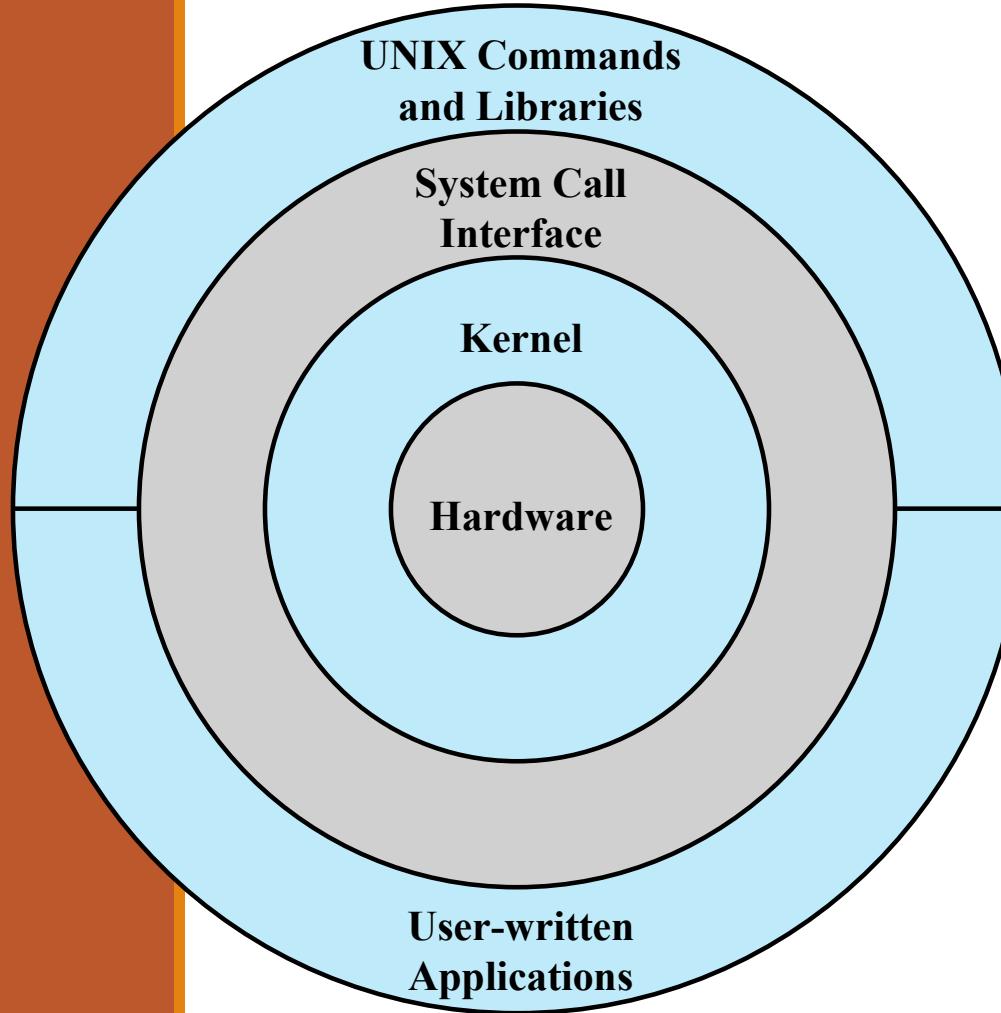


Figure 2.15 General UNIX Architecture

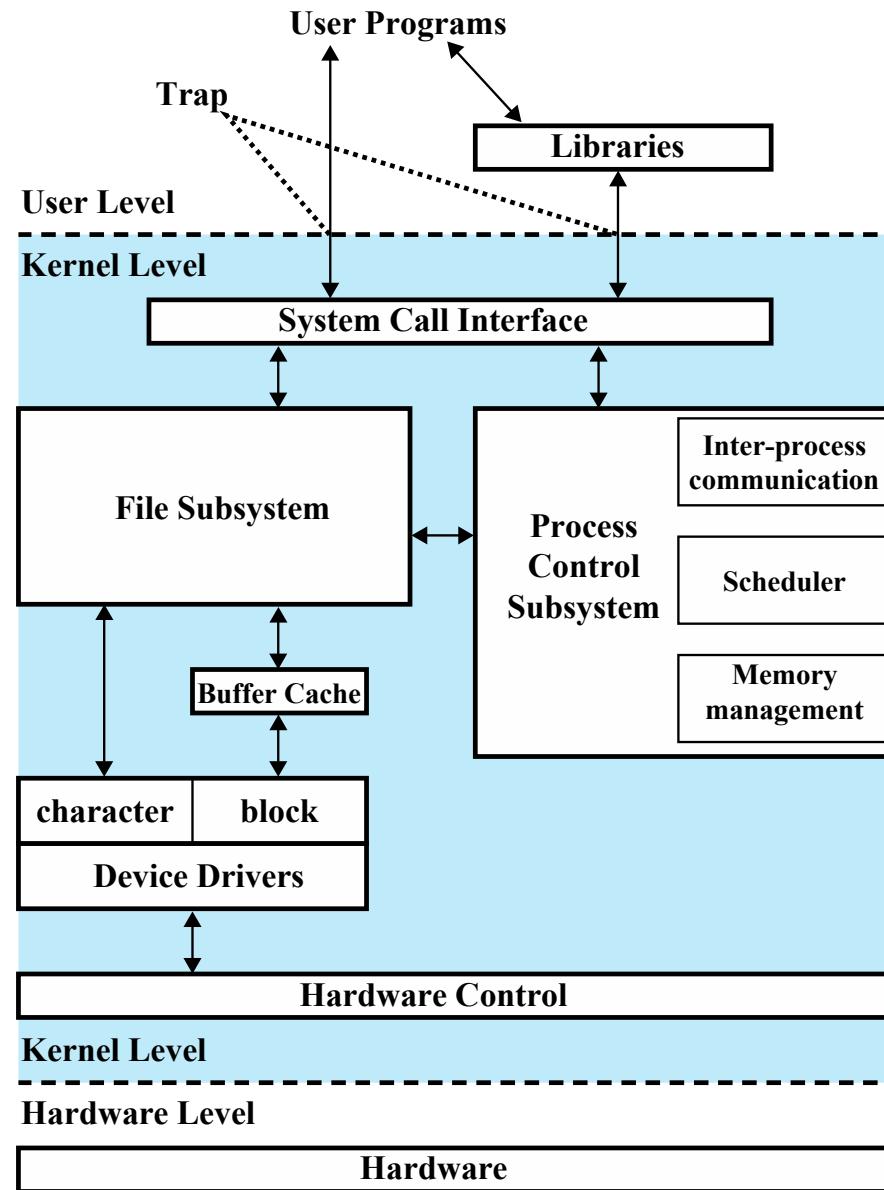


Figure 2.16 Traditional UNIX Kernel

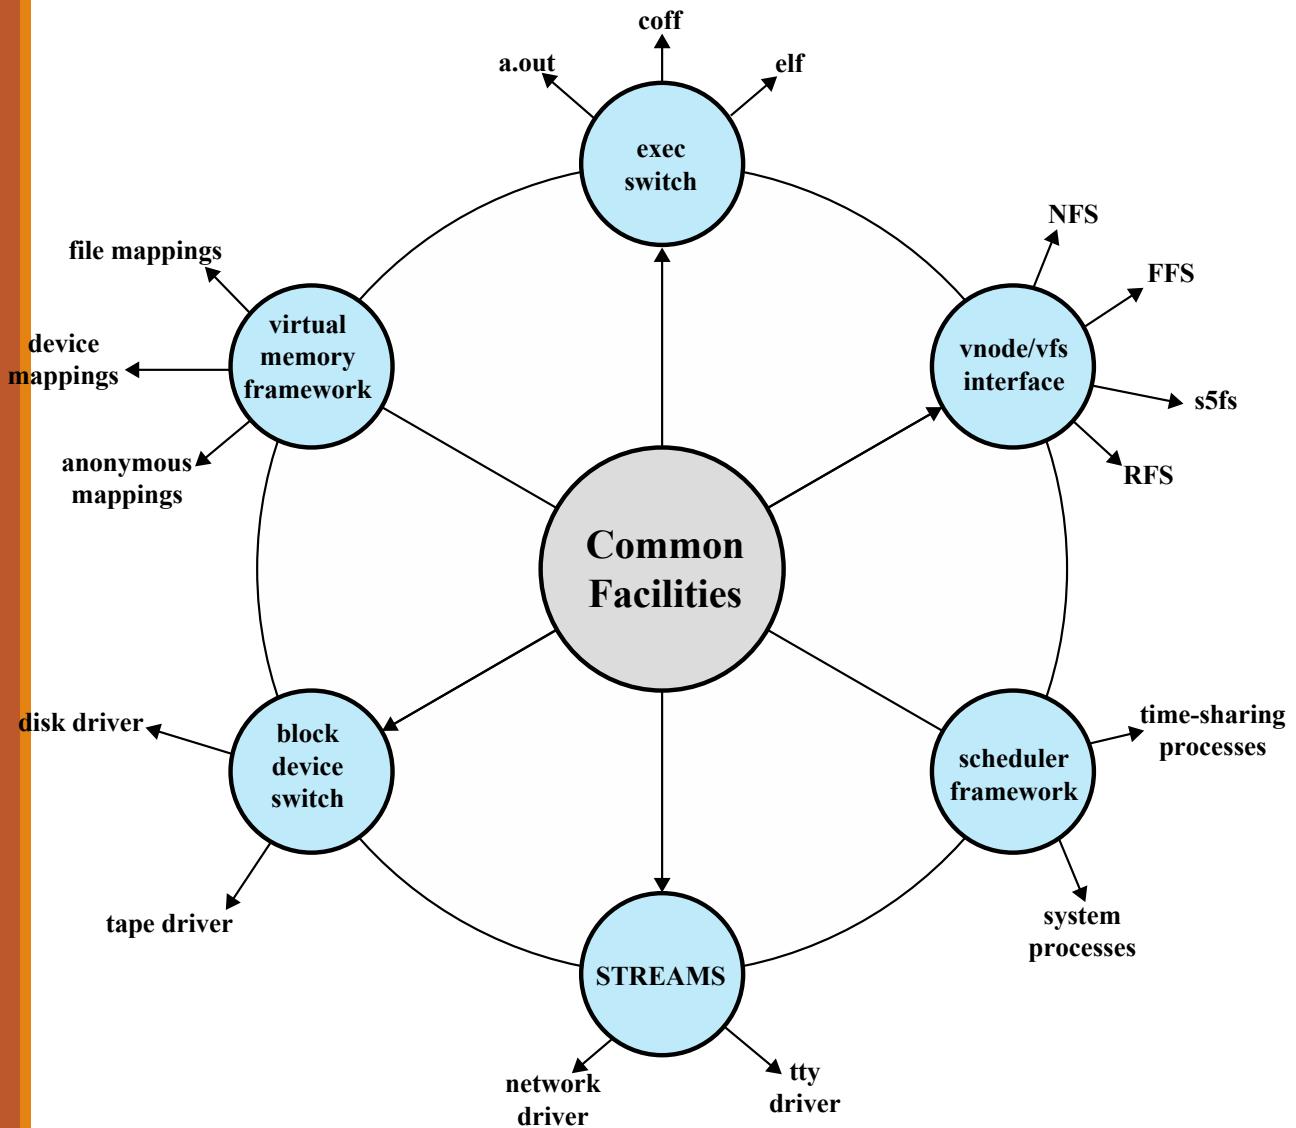


Figure 2.17 Modern UNIX Kernel [VAHA96]

System V Release 4 (SVR4)

Developed jointly by AT&T and Sun Microsystems

Combines features from SVR3, 4.3BSD, Microsoft Xenix
System V, and SunOS

New features in the release include:

- real-time processing support
- process scheduling classes
- dynamically allocated data structures
- virtual memory management
- virtual file system
- preemptive kernel

BSD

Berkeley Software Distribution

4.xBSD is widely used in academic installations and has served as the basis of a number of commercial UNIX products

4.4BSD was the final version of BSD to be released by Berkeley

- major upgrade to 4.3BSD
- includes
 - a new virtual memory system
 - changes in the kernel structure
 - several other feature enhancements

FreeBSD

- one of the most widely used and best documented versions
- popular for Internet-based servers and firewalls
- used in a number of embedded systems
- Mac OS X is based on FreeBSD 5.0 and the Mach 3.0 microkernel

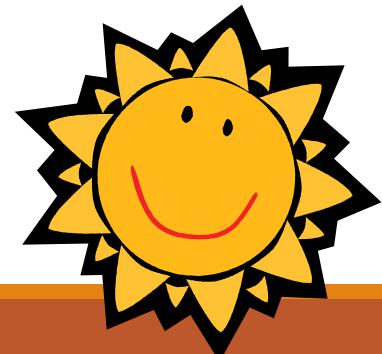
Solaris 10

Sun's SVR4-based UNIX release

Provides all of the features of SVR4 plus a number of more advanced features such as:

- a fully preemptable, multithreaded kernel
- full support for SMP
- an object-oriented interface to file systems

Most widely used and most successful commercial UNIX implementation



LINUX Overview

Started out as a UNIX variant for the IBM PC

Linus Torvalds, a Finnish student of computer science, wrote the initial version

Linux was first posted on the Internet in 1991

Today it is a full-featured UNIX system that runs on several platforms

Is free and the source code is available

Key to success has been the availability of free software packages

Highly modular and easily configured

Modular Monolithic Kernel

- INCLUDES VIRTUALLY ALL OF THE OS FUNCTIONALITY IN ONE LARGE BLOCK OF CODE THAT RUNS AS A SINGLE PROCESS WITH A SINGLE ADDRESS SPACE
- ALL THE FUNCTIONAL COMPONENTS OF THE KERNEL HAVE ACCESS TO ALL OF ITS INTERNAL DATA STRUCTURES AND ROUTINES
- LINUX IS STRUCTURED AS A COLLECTION OF MODULES

LOADABLE MODULES

Relatively independent blocks

A module is an object file whose code can be linked to and unlinked from the kernel at runtime

A module is executed in kernel mode on behalf of the current process

Have two important characteristics:

- dynamic linking
- stackable modules

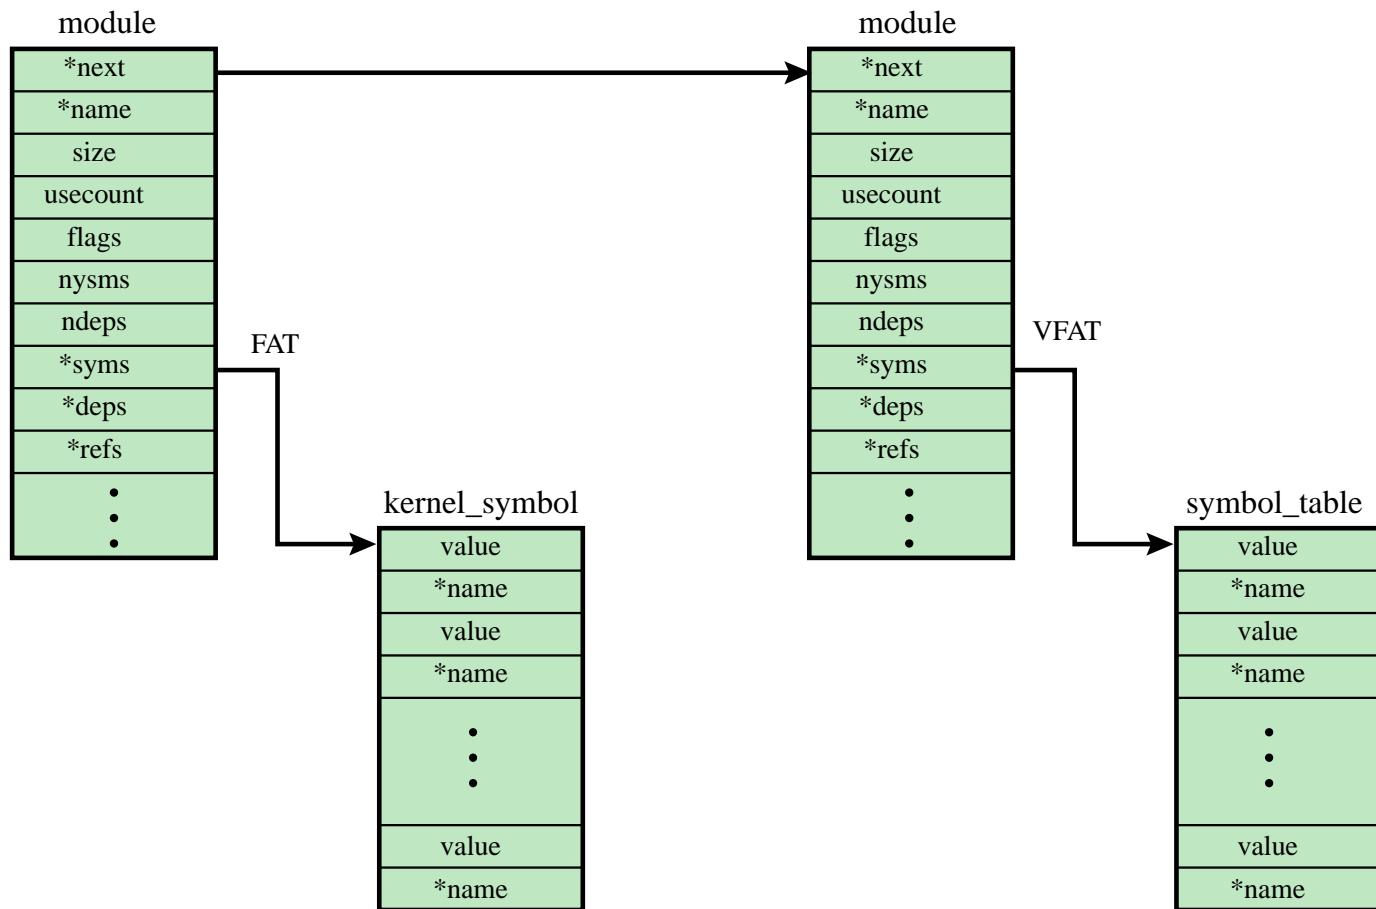


Figure 2.18 Example List of Linux Kernel Modules

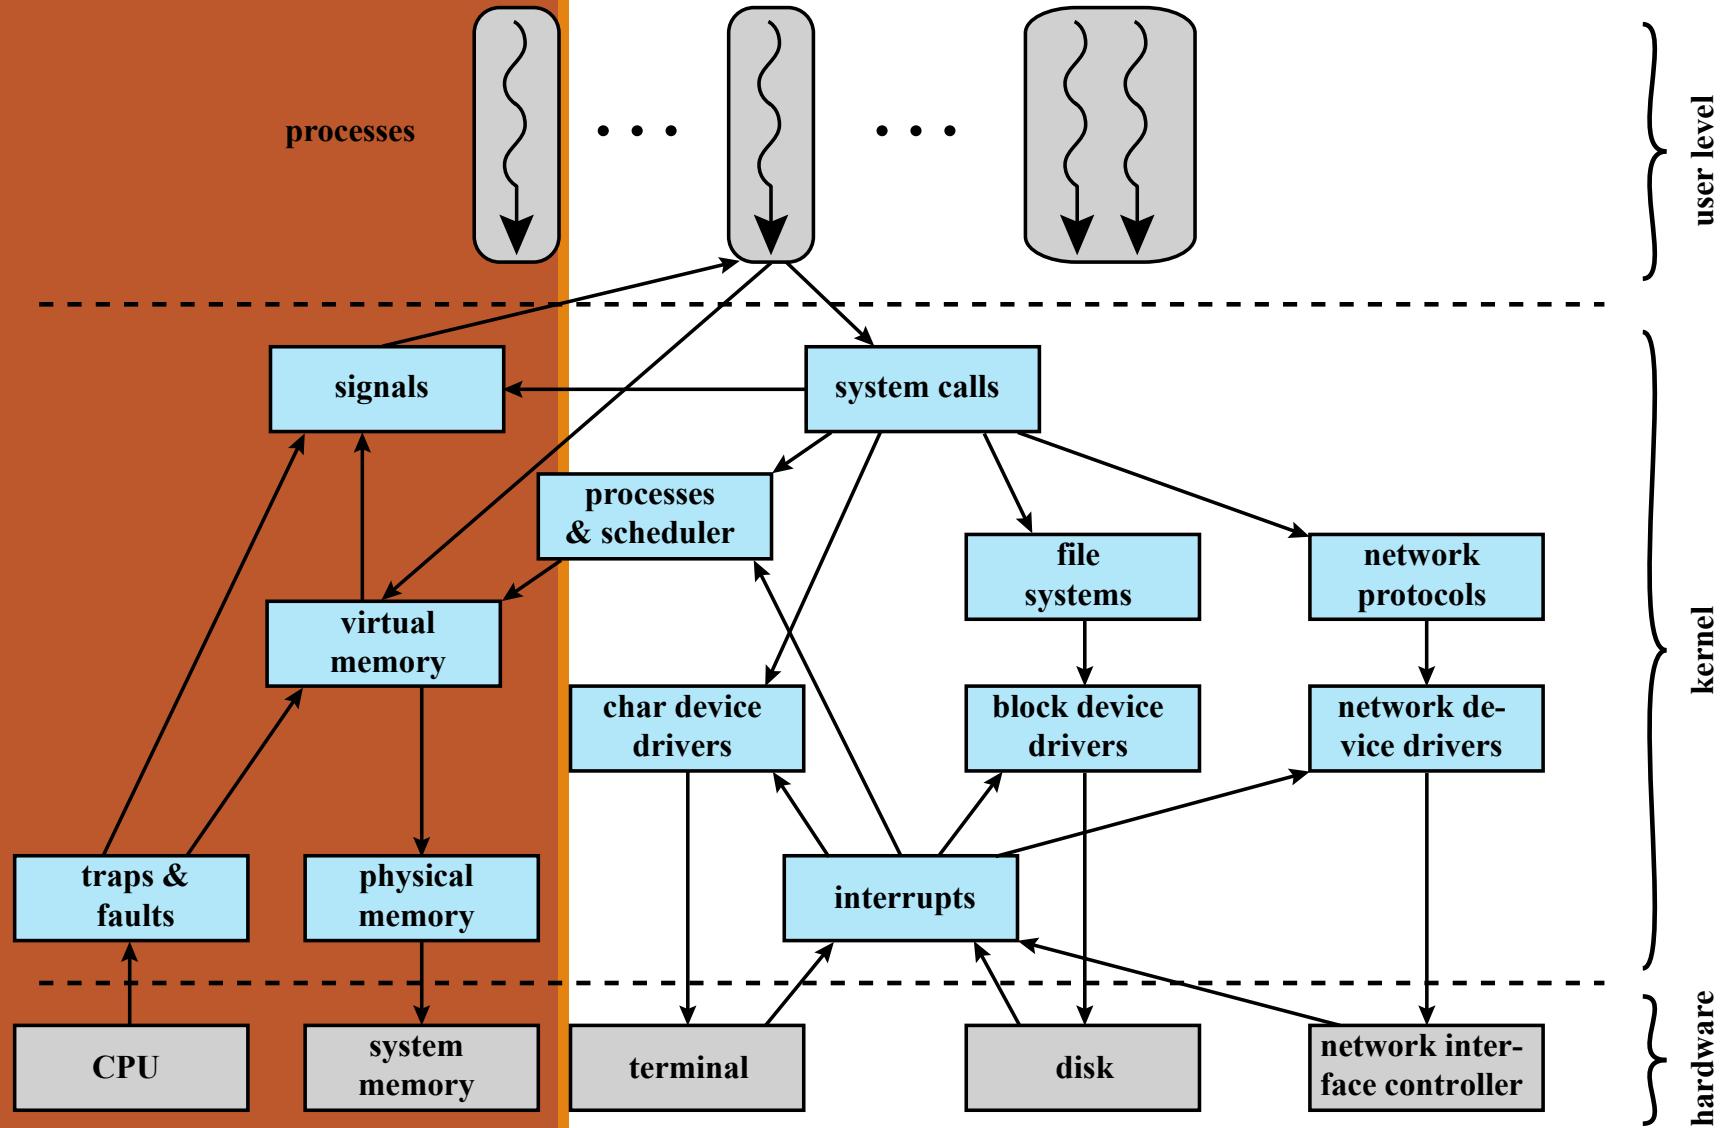


Figure 2.19 Linux Kernel Components

Linux Signals

SIGHUP	Terminal hangup	SIGCONT	Continue
SIGQUIT	Keyboard quit	SIGTSTP	Keyboard stop
SIGTRAP	Trace trap	SIGTTOU	Terminal write
SIGBUS	Bus error	SIGXCPU	CPU limit exceeded
SIGKILL	Kill signal	SIGVTALRM	Virtual alarm clock
SIGSEGV	Segmentation violation	SIGWINCH	Window size unchanged
SIGPIPT	Broken pipe	SIGPWR	Power failure
SIGTERM	Termination	SIGRTMIN	First real-time signal
SIGCHLD	Child status unchanged	SIGRTMAX	Last real-time signal

Table 2.6 Some Linux Signals

Filesystem related	
close	Close a file descriptor.
link	Make a new name for a file.
open	Open and possibly create a file or device.
read	Read from file descriptor.
write	Write to file descriptor
Process related	
execve	Execute program.
exit	Terminate the calling process.
getpid	Get process identification.
setuid	Set user identity of the current process.
ptrace	Provides a means by which a parent process may observe and control the execution of another process, and examine and change its core image and registers.
Scheduling related	
sched_getparam	Sets the scheduling parameters associated with the scheduling policy for the process identified by <code>pid</code> .
sched_get_priority_max	Returns the maximum priority value that can be used with the scheduling algorithm identified by <code>policy</code> .
sched_setscheduler	Sets both the scheduling policy (e.g., FIFO) and the associated parameters for the process <code>pid</code> .
sched_rr_get_interval	Writes into the timespec structure pointed to by the parameter <code>tp</code> the round robin time quantum for the process <code>pid</code> .
sched_yield	A process can relinquish the processor voluntarily without blocking via this system call. The process will then be moved to the end of the queue for its static priority and a new process gets to run.

Table 2.7 Some Linux System Calls (page 1 of 2)

Interprocess Communication (IPC) related	
msgrecv	A message buffer structure is allocated to receive a message. The system call then reads a message from the message queue specified by msqid into the newly created message buffer.
semctl	Performs the control operation specified by cmd on the semaphore set semid.
semop	Performs operations on selected members of the semaphore set semid.
shmat	Attaches the shared memory segment identified by shmid to the data segment of the calling process.
shmctl	Allows the user to receive information on a shared memory segment, set the owner, group, and permissions of a shared memory segment, or destroy a segment.
Socket (networking) related	
bind	Assigns the local IP address and port for a socket. Returns 0 for success and -1 for error.
connect	Establishes a connection between the given socket and the remote socket associated with sockaddr.
gethostname	Returns local host name.
send	Send the bytes contained in buffer pointed to by *msg over the given socket.
setsockopt	Sets the options on a socket
Miscellaneous	
fsync	Copies all in-core parts of a file to disk, and waits until the device reports that all parts are on stable storage.
time	Returns the time in seconds since January 1, 1970.
vhangup	Simulates a hangup on the current terminal. This call arranges for other users to have a "clean" tty at login time.

Table 2.7 Some Linux System Calls (page 2 of 2)

Android Operating System

A Linux-based system originally designed for touchscreen mobile devices such as smartphones and tablet computers

The most popular mobile OS

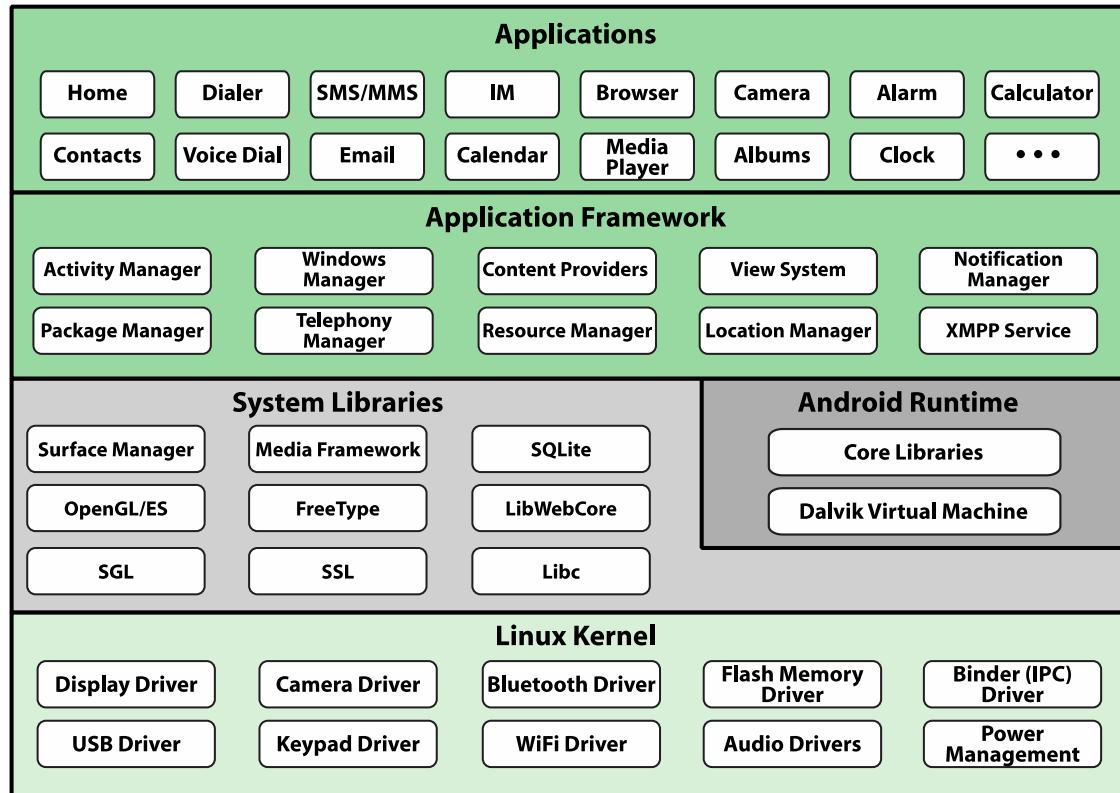
Development was done by Android Inc., which was bought by Google in 2005

1st commercial version (Android 1.0) was released in 2008

Most recent version is Android 4.3 (Jelly Bean)

The Open Handset Alliance (OHA) was responsible for the Android OS releases as an open platform

The open-source nature of Android has been the key to its success



Implementation:

Applications, Application Framework: Java

System Libraries, Android Runtime: C and C++

Linux Kernel: C

Figure 2.20 Android Software Architecture

Application Framework

Provides high-level building blocks accessible through standardized API's that programmers use to create new apps

- architecture is designed to simplify the reuse of components

Key components:

Activity Manager

Manages lifecycle of applications

Responsible for starting, stopping, and resuming the various applications

Window Manager

Java abstraction of the underlying Surface Manager

Allows applications to declare their client area and use features like the status bar

Package Manager

Installs and removes applications

Telephony Manager

Allows interaction with phone, SMS, and MMS services

Application Framework

(cont.)

- Key components: (cont.)
 - Content Providers
 - these functions encapsulate application data that need to be shared between applications such as contacts
 - Resource Manager
 - manages application resources, such as localized strings and bitmaps
 - View System
 - provides the user interface (UI) primitives as well as UI Events
 - Location Manager
 - allows developers to tap into location-based services, whether by GPS, cell tower IDs, or local Wi-Fi databases
 - Notification Manager
 - manages events, such as arriving messages and appointments
 - XMPP
 - provides standardized messaging functions between applications

System Libraries

Collection of useful system functions written in C or C++ and used by various components of the Android system

Called from the application framework and applications through a Java interface

Exposed to developers through the Android application framework

Some of the key system libraries include:

- Surface Manager
- OpenGL
- Media Framework
- SQL Database
- Browser Engine
- Bionic LibC

- Every Android application runs in its own process with its own instance of the Dalvik virtual machine (DVM)
- DVM executes files in the Dalvik Executable (.dex) format
- Component includes a set of core libraries that provides most of the functionality available in the core libraries of the Java programming language
- To execute an operation the DVM calls on the corresponding C/C++ library using the Java Native Interface (JNI)

Android Runtime



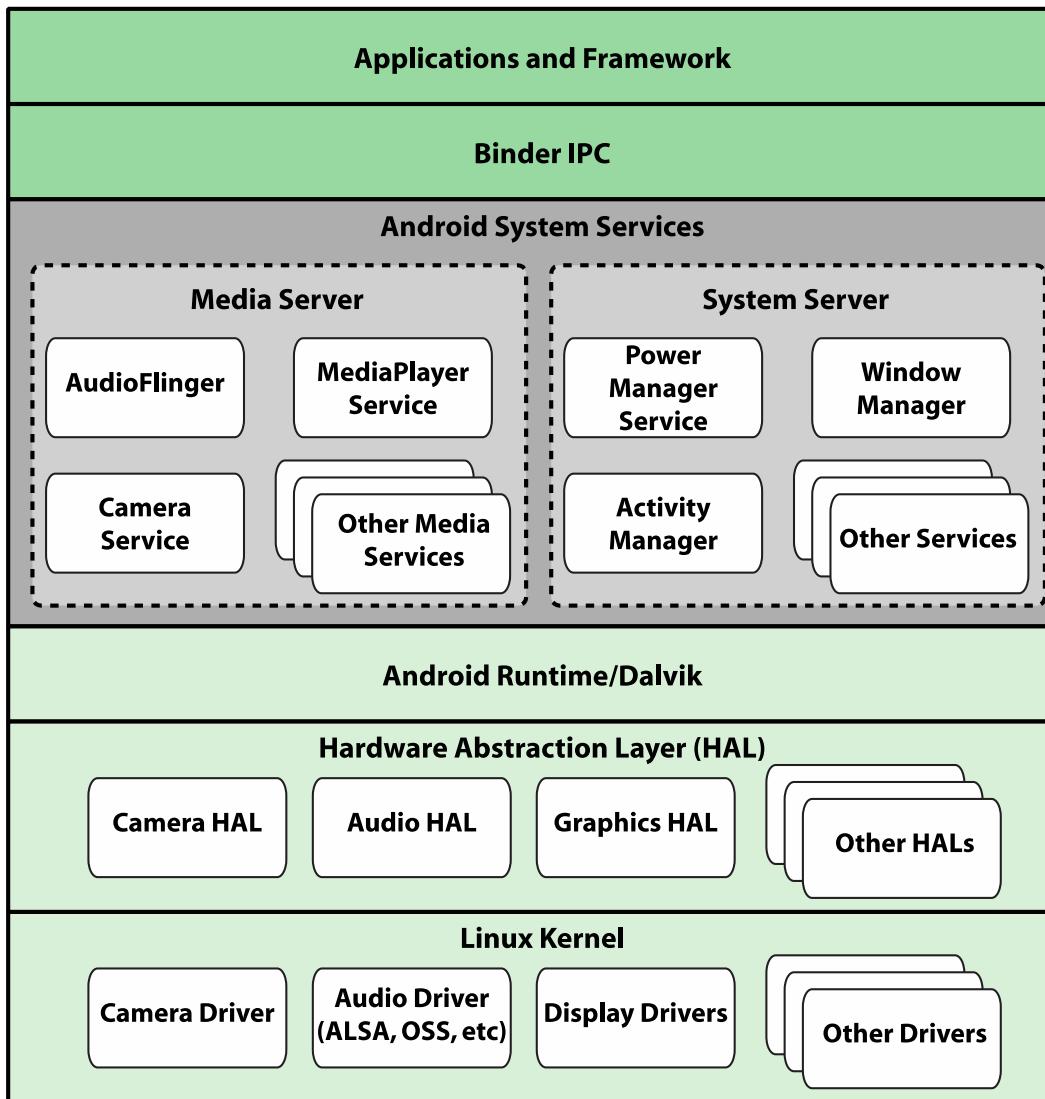


Figure 2.21 Android System Architecture

Activities

An activity is a single visual user interface component, including things such as menu selections, icons, and checkboxes

Every screen in an application is an extension of the Activity class

Use Views to form graphical user interfaces that display information and respond to user actions

Power Management

ALARMS

Implemented in the Linux kernel and is visible to the app developer through the AlarmManager in the RunTime core libraries

Is implemented in the kernel so that an alarm can trigger even if the system is in sleep mode

- this allows the system to go into sleep mode, saving power, even though there is a process that requires a wake up

WAKELOCKS

Prevents an Android system from entering into sleep mode

These locks are requested through the API whenever an application requires one of the managed peripherals to remain powered on

An application can hold one of the following wakelocks:

- Full_Wake_Lock
- Partial_Wake_Lock
- Screen_Dim_Wake_Lock
- Screen_Bright_Wake_Lock

Summary

Operating system objectives and functions

- User/computer interface
- Resource manager
- Evolution of operating systems
- Serial processing
- Simple/multiprogrammed/time-sharing batch systems
- Major achievements
- Developments leading to modern operating systems
- Fault tolerance
- Fundamental concepts
- Faults
- OS mechanisms

- OS design considerations for multiprocessor and multicore

- Microsoft Windows overview

Traditional Unix systems

- History/description

Modern Unix systems

- System V Release 4 (SVR4)
- BSD
- Solaris 10

Linux

- History
- Modular structure
- Kernel components

Android

- Software/system architecture
- Activities

Power management